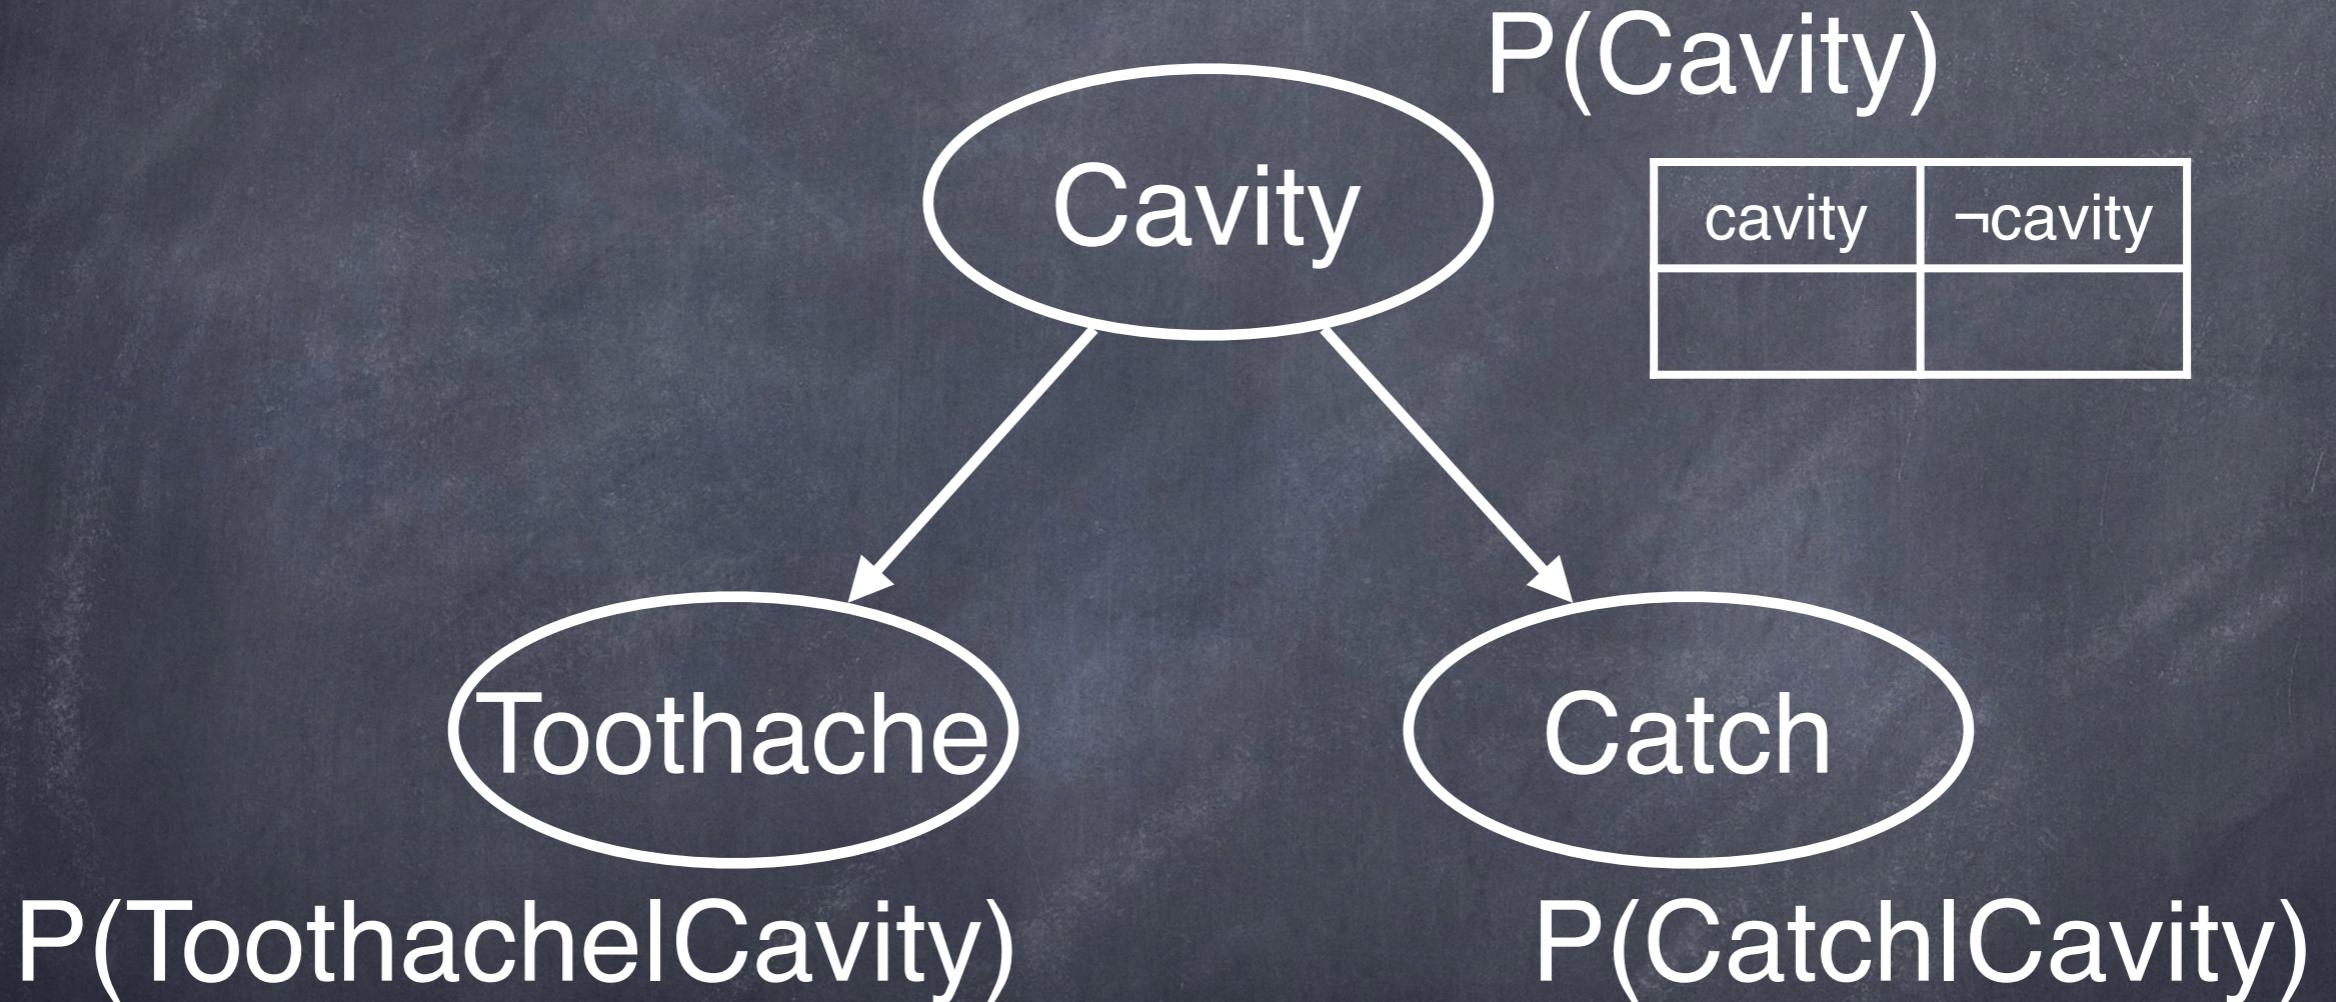


CSC442: Introduction to Artificial Intelligence

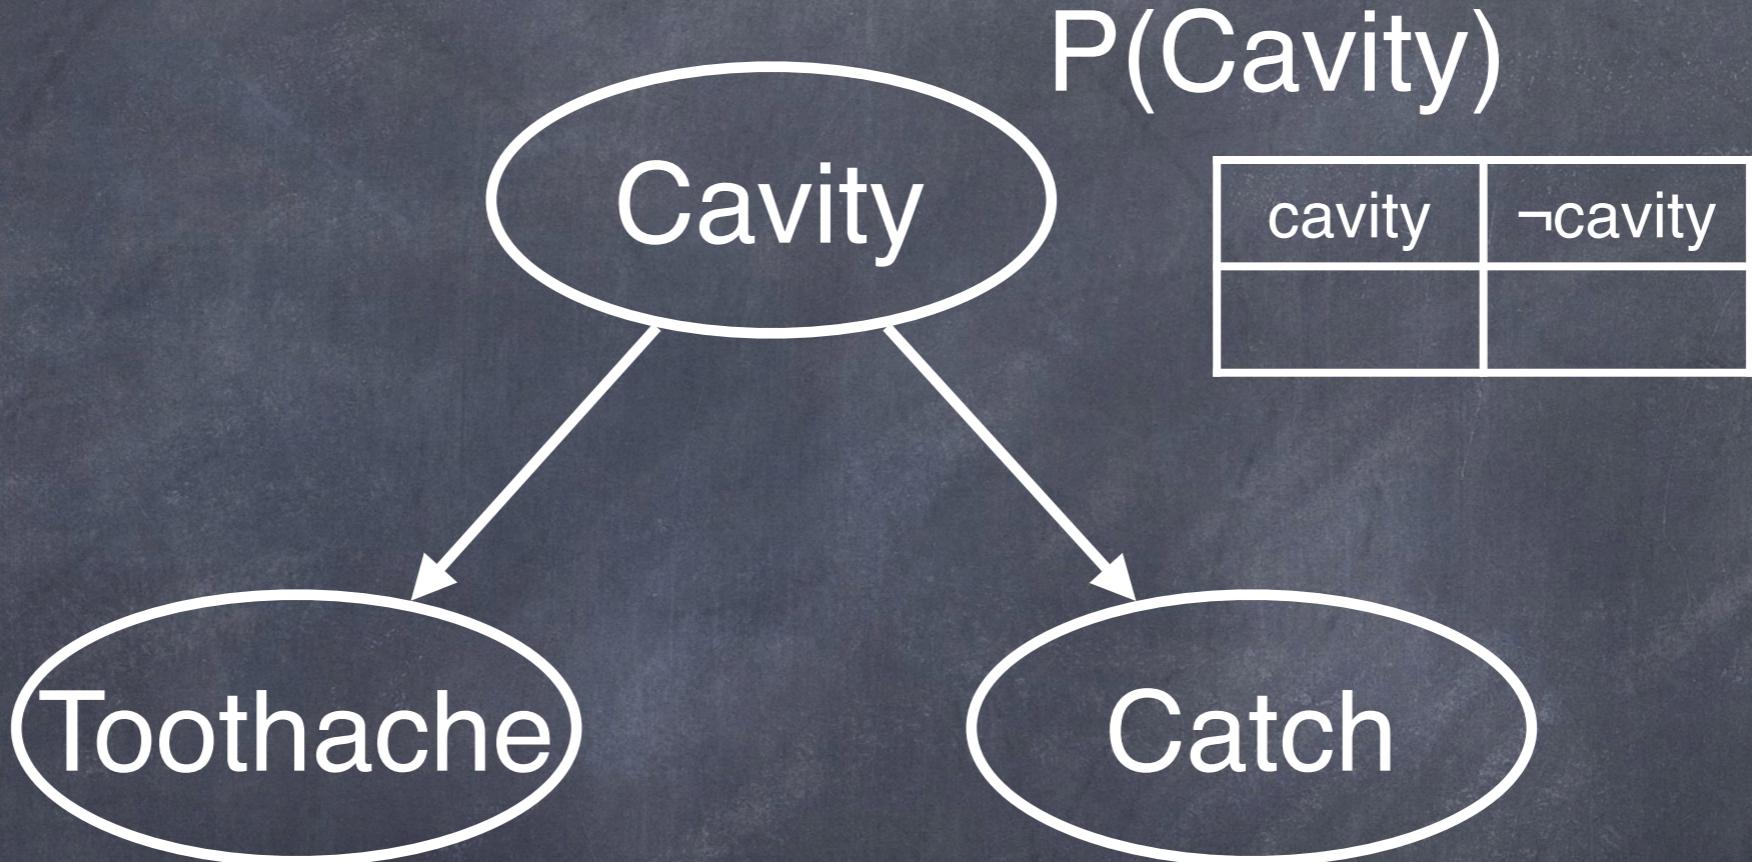
Lecture 3.4

Bayesian Networks



Cavity	toothache	$\neg\text{toothache}$
cavity		
$\neg\text{cavity}$		

Cavity	catch	$\neg\text{catch}$
cavity		
$\neg\text{cavity}$		

$P(\text{toothache}, \text{cavity}, \text{catch}) =$ $P(\text{toothache}|\text{cavity})P(\text{catch}|\text{cavity})P(\text{cavity})$  $P(\text{Cavity})$

cavity	\neg cavity

 $P(\text{Toothache}|\text{Cavity})$

Cavity	toothache	\neg toothache
cavity		
\neg cavity		

 $P(\text{Catch}|\text{Cavity})$

Cavity	catch	\neg catch
cavity		
\neg cavity		

$P(B)$

$P(b)$
0.001

Burglary

Alarm

Earthquak

$P(E)$

$P(e)$
0.002

$P(A|B, E)$

B	E	$P(a B, E)$
t	t	0.95
t	f	0.94
f	t	0.29
f	f	0.001

$P(J|A)$

A	$P(j A)$
t	0.9
f	0.05

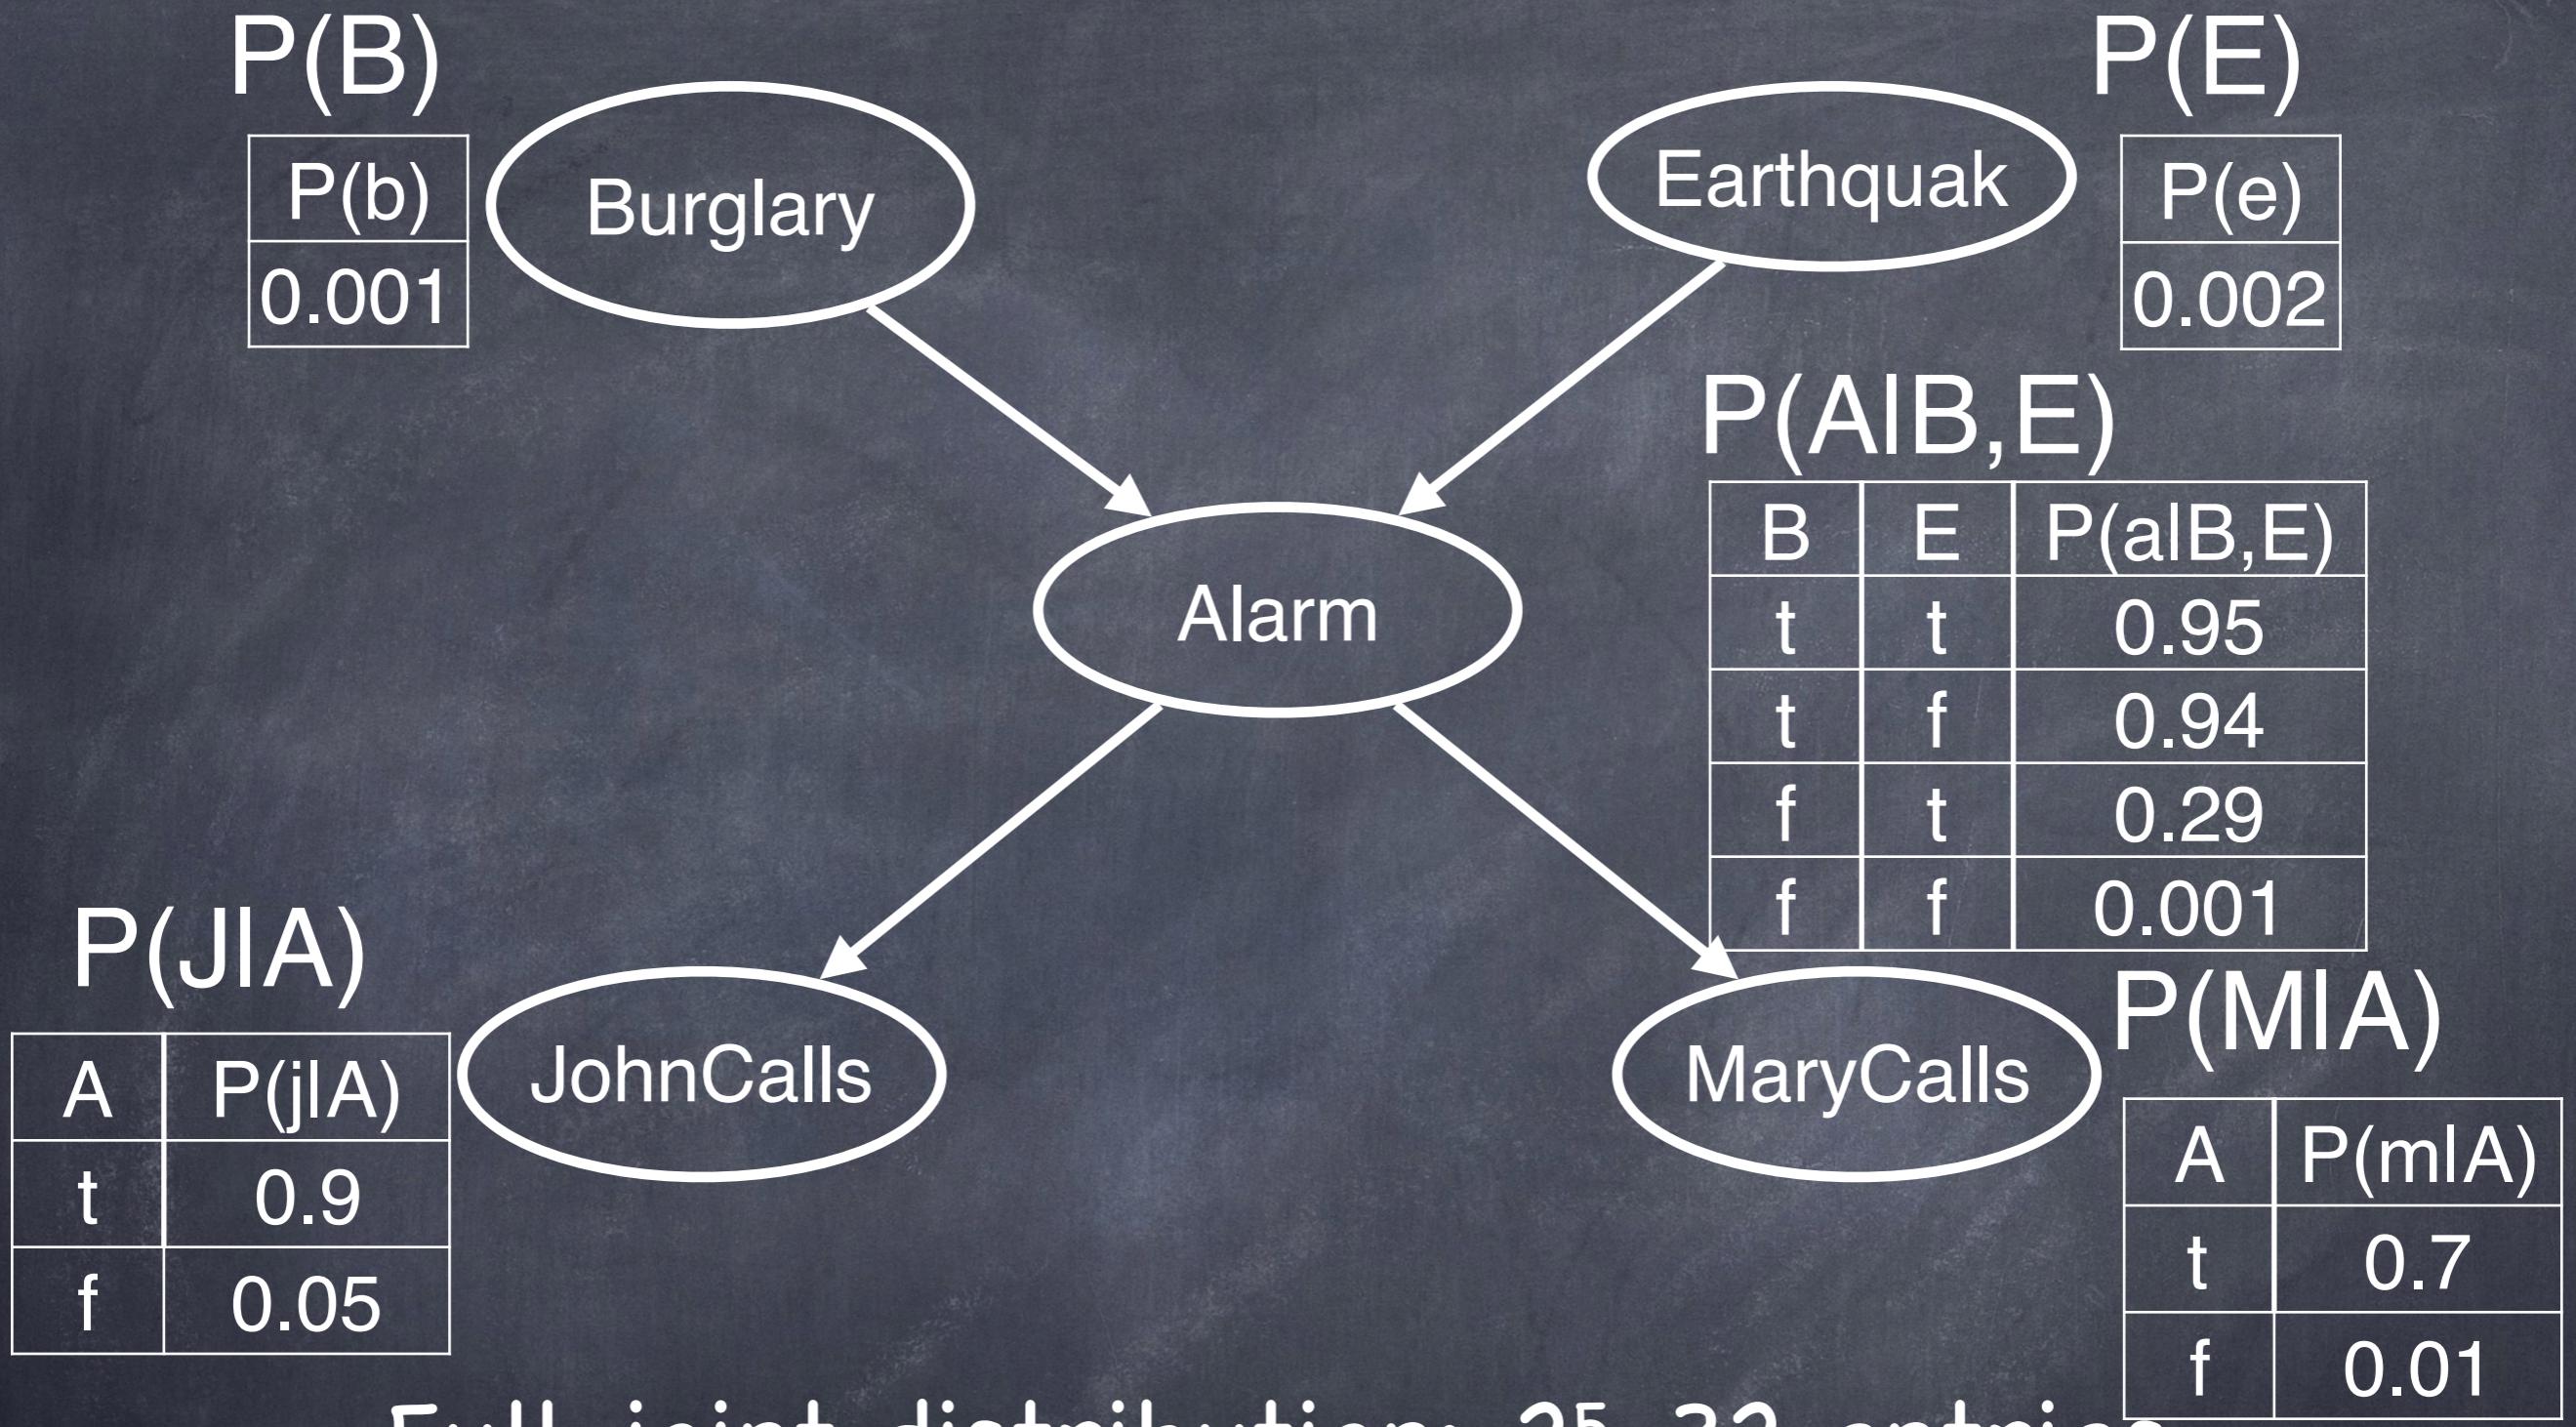
JohnCalls

MaryCalls

$P(M|A)$

A	$P(m A)$
t	0.7
f	0.01

$$P(B, E, A, J, M) = \alpha P(B) P(E) P(A | B, E) P(J | A) P(M | A)$$



Full joint distribution: $2^5=32$ entries

Bayesian network: 10 entries

Assuming conditional independences
encoded in the network

Inference in Bayesian Networks

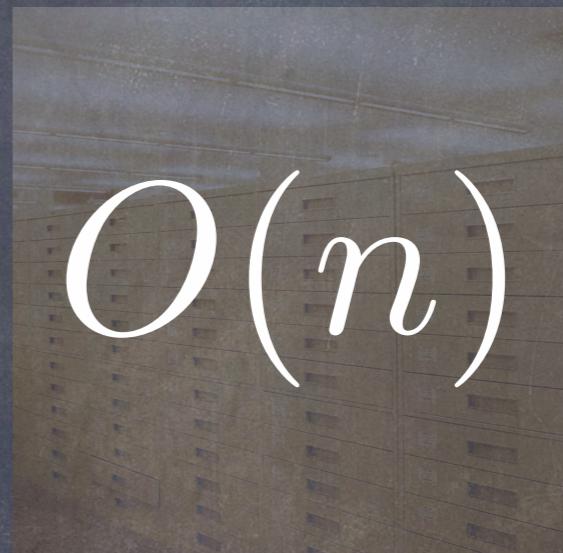
$$\begin{aligned} \mathbf{P}(X \mid \mathbf{e}) &= \alpha \mathbf{P}(X, \mathbf{e}) = \alpha \sum_{\mathbf{y}} \mathbf{P}(X, \mathbf{e}, \mathbf{y}) \\ &= \alpha \sum_{\mathbf{y}} \prod_{i=1}^n P(X_i \mid \text{parents}(X_i)) \end{aligned}$$

- “A query can be answered from a Bayesian Network by computing sums of products of conditional probabilities from the network.”

Exact Inference in BNs



Time Complexity



Space Complexity

Approximate Inference in Bayesian Networks

Exact Inference

- Query variable X
- Evidence variables E_1, \dots, E_m
 - Observed values: $e = \langle e_1, \dots, e_m \rangle$
- Non-evidence, non-query (“hidden”) variables: Y
- Compute: $P(X | e)$

Approximate Inference

- Query variable X
- Evidence variables E_1, \dots, E_m
 - Observed values: $e = \langle e_1, \dots, e_m \rangle$
- Non-evidence, non-query (“hidden”) variables: Y
- Approximate (estimate): $P(X | e)$

Unconditional Approximation

- Query variable X
- Evidence variables E_1, \dots, E_m
 - Observed values: $e = \langle e_1, \dots, e_m \rangle$
- Non-evidence, non-query (“hidden”) variables: Y
- Approximate (estimate): $P(X)$

Heads

Heads

Goal: $P(\text{Heads})$
 $P(\text{Heads}=\text{true})$
 $P(\text{heads})$

Heads

of flips: N

of heads: N_{heads}

Heads

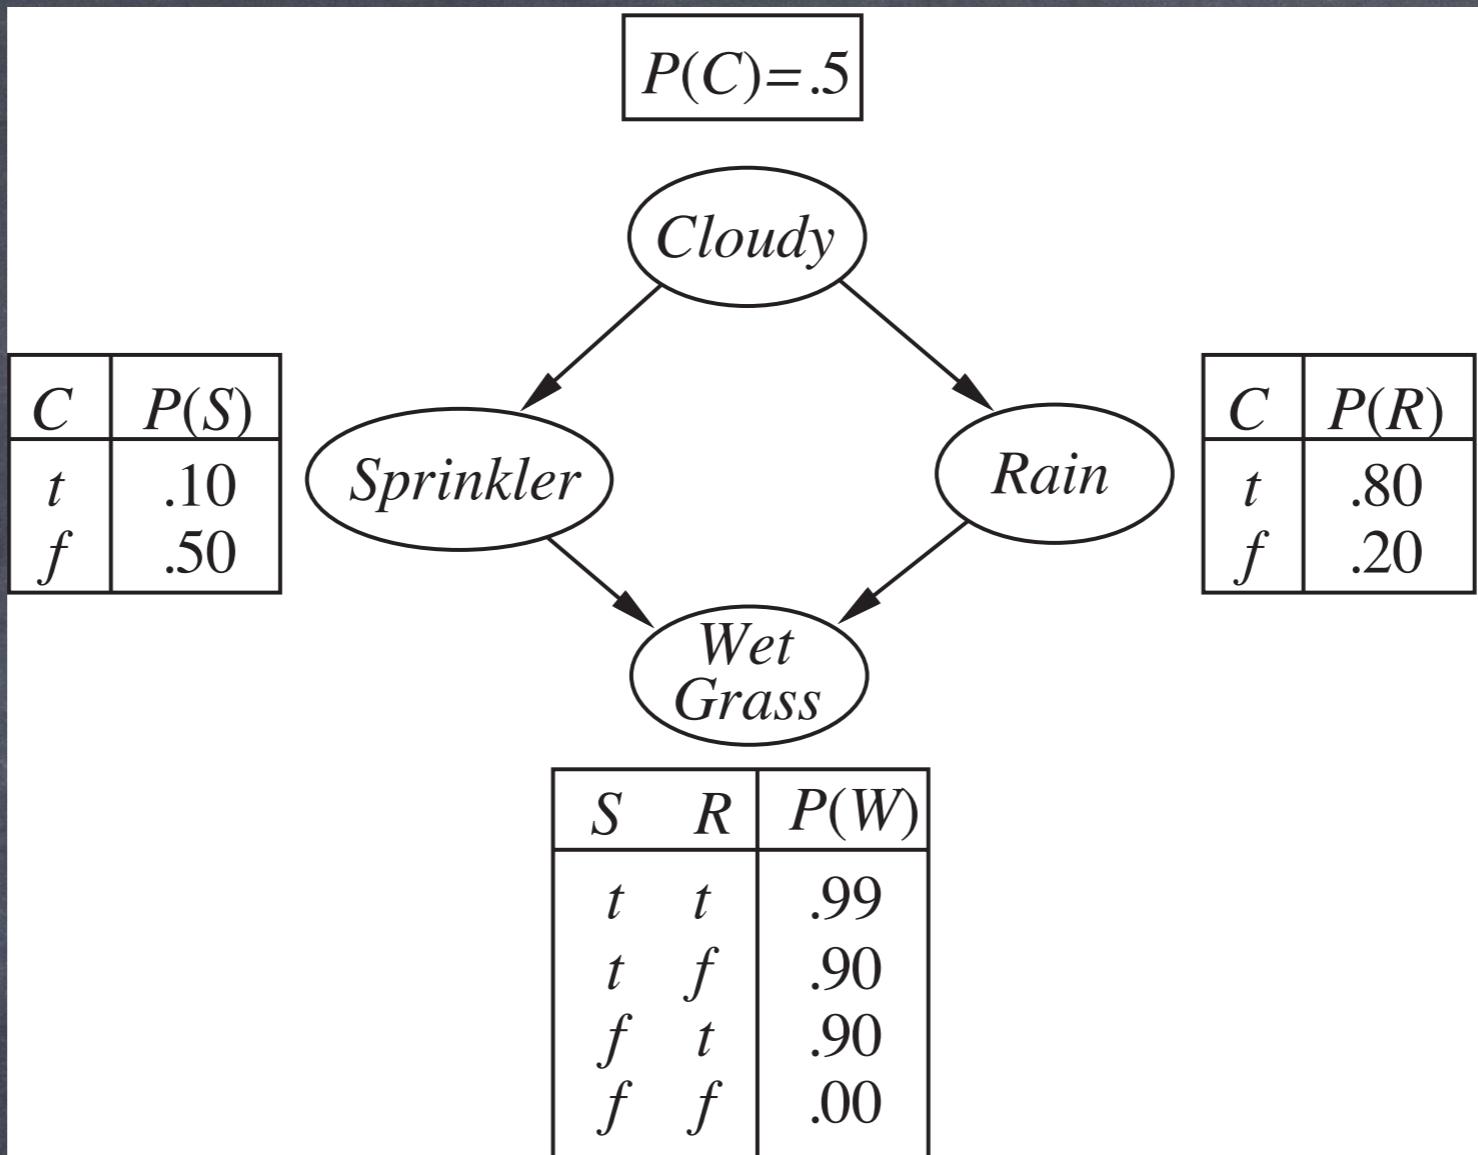
$$P(\text{heads}) \approx \frac{N_{\text{heads}}}{N}$$

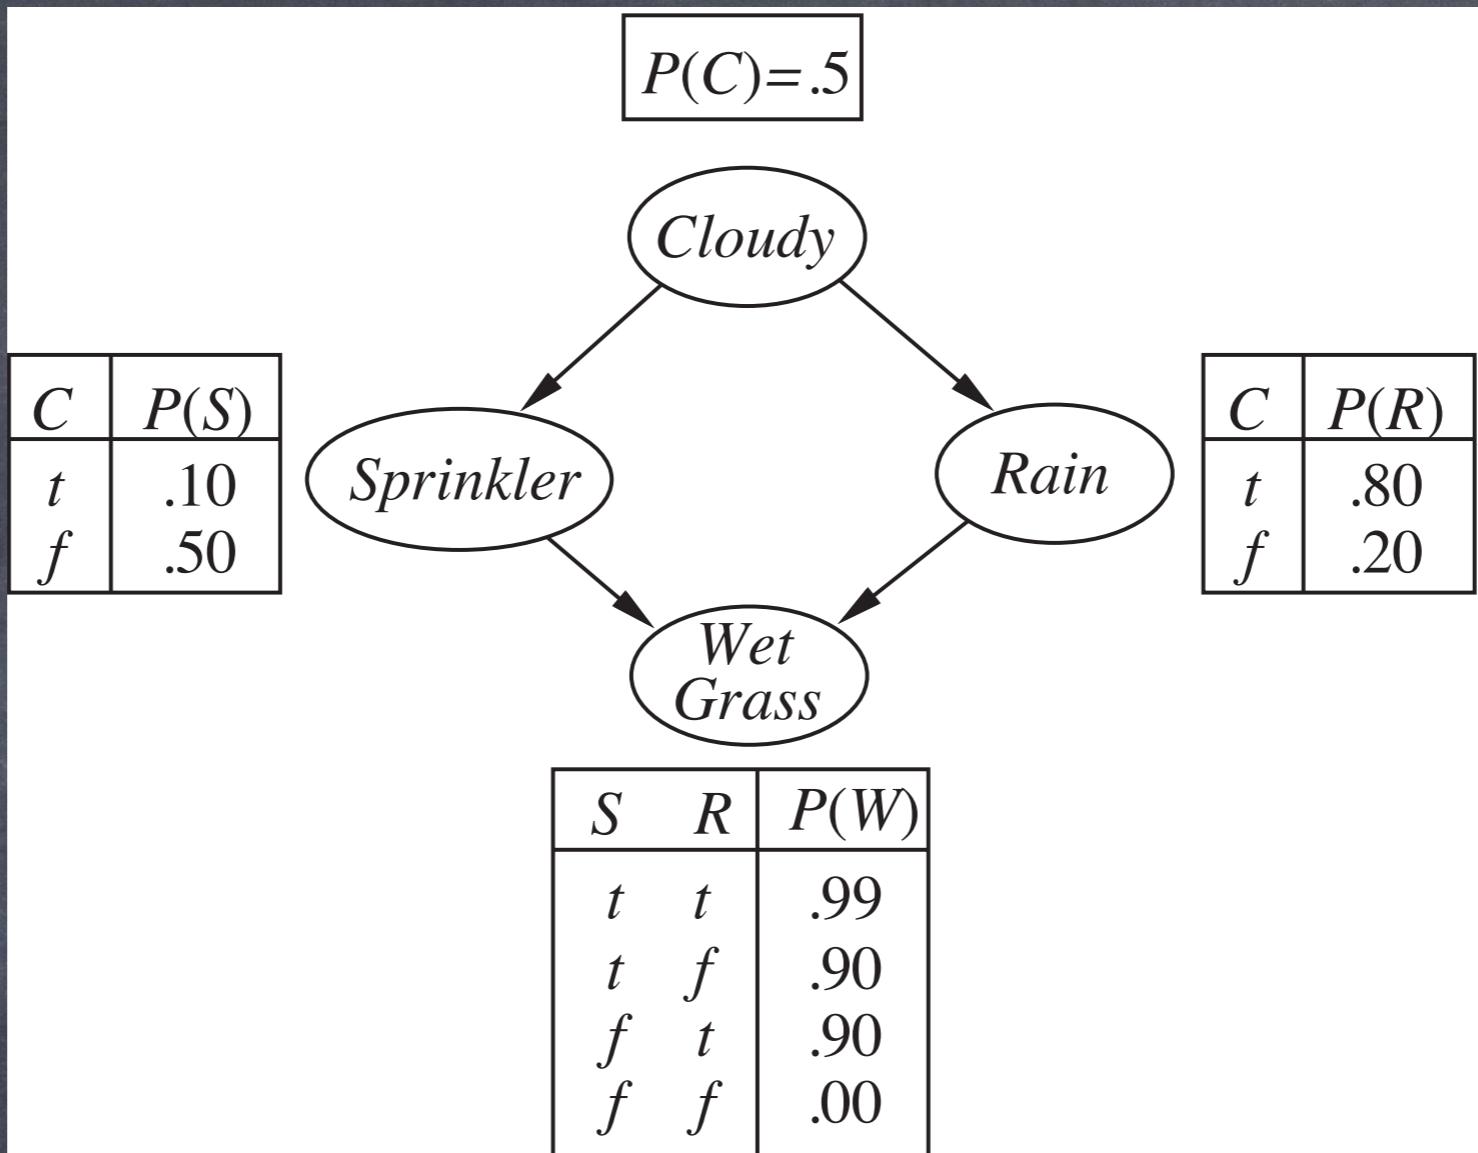
Heads

$$P(\text{heads}) = \lim_{N \rightarrow \infty} \frac{N_{\text{heads}}}{N}$$

Sampling

- Generating events (possible worlds) from a distribution
- Estimating probabilities as ratio of observed events to total events
- Consistent estimate: becomes exact in the large-sample limit

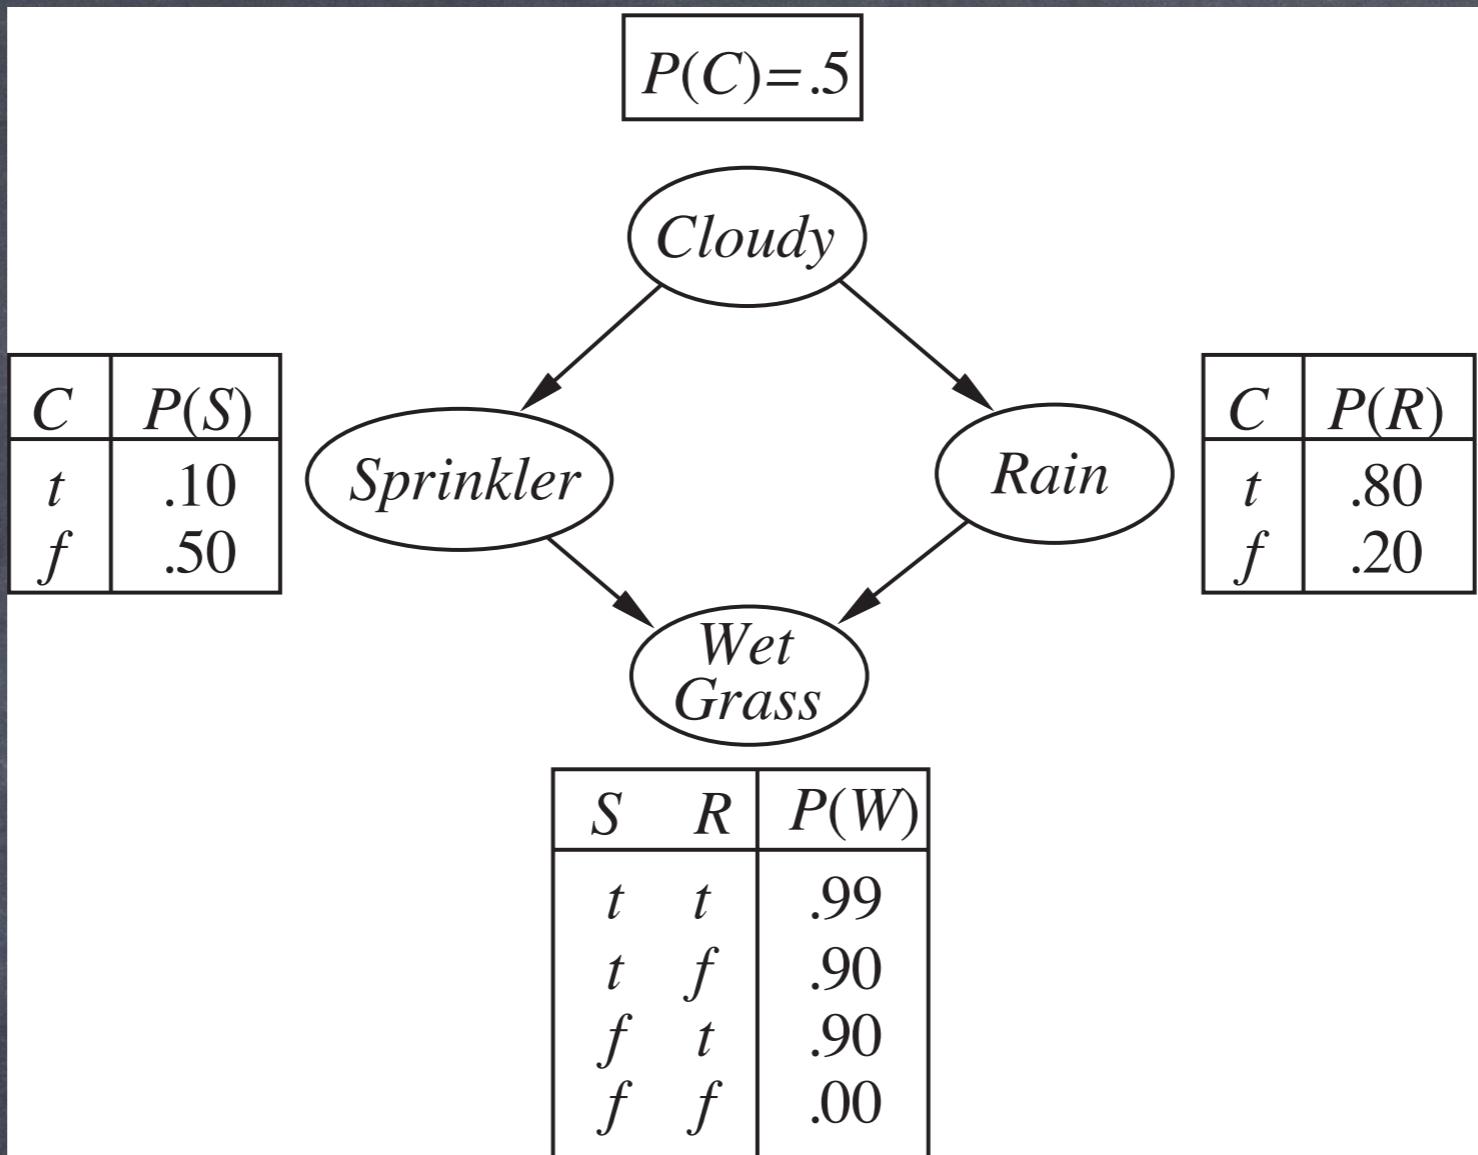




$$P(\text{Rain} = \text{true})$$

Sampling

- Generate assignments of values to the random variables
- That are consistent with the full joint distribution encoded in the network
 - In the sense that in the limit, the frequency of occurrence of any event (possible world) is equal to its probability



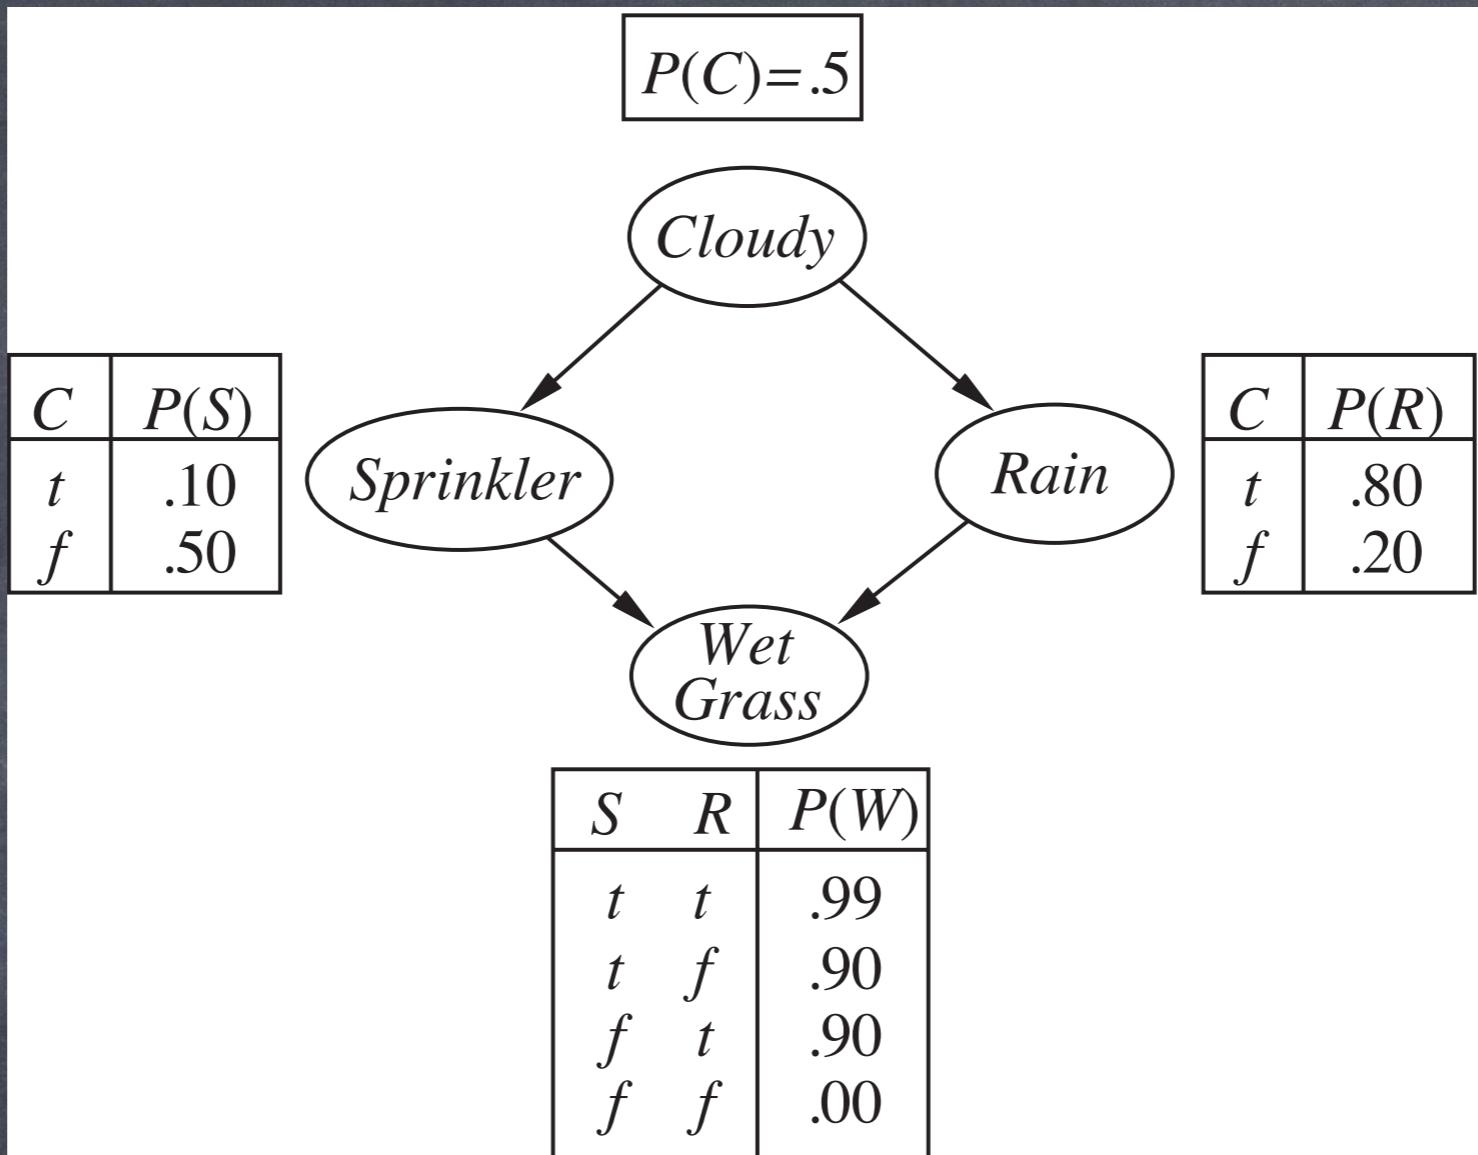
$\langle \text{Cloudy}=\text{true}, \text{Sprinkler}=\text{false}, \text{Rain}=\text{true}, \text{WetGrass}=\text{true} \rangle$

likelihood
weighing :

$$P(R | +C, +W)$$

$$W = 1.0$$

- ① $C \rightarrow \text{evidence}$ $+C$
 ② $S \rightarrow \text{non-evidence}$ $-S$



$\langle \text{Cloudy}=\text{true}, \text{Sprinkler}=\text{false}, \text{Rain}=\text{true}, \text{WetGrass}=\text{true} \rangle$

$\langle \text{Cloudy}=\text{false}, \text{Sprinkler}=\text{false}, \text{Rain}=\text{false}, \text{WetGrass}=\text{false} \rangle$

$\langle \text{Cloudy}=\text{true}, \text{Sprinkler}=\text{true}, \text{Rain}=\text{true}, \text{WetGrass}=\text{true} \rangle$

...

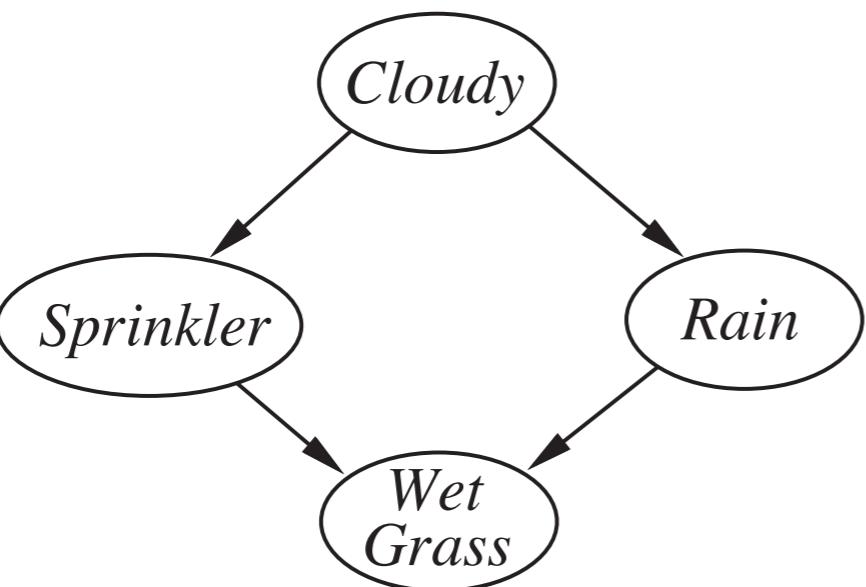


Generating Samples

- Sample each variable in topological order
 - Child appears after its parents
 - Choose the value for that variable conditioned on the values already chosen for its parents

Cloudy
Sprinkler
Rain
WetGrass

$$P(C) = .5$$



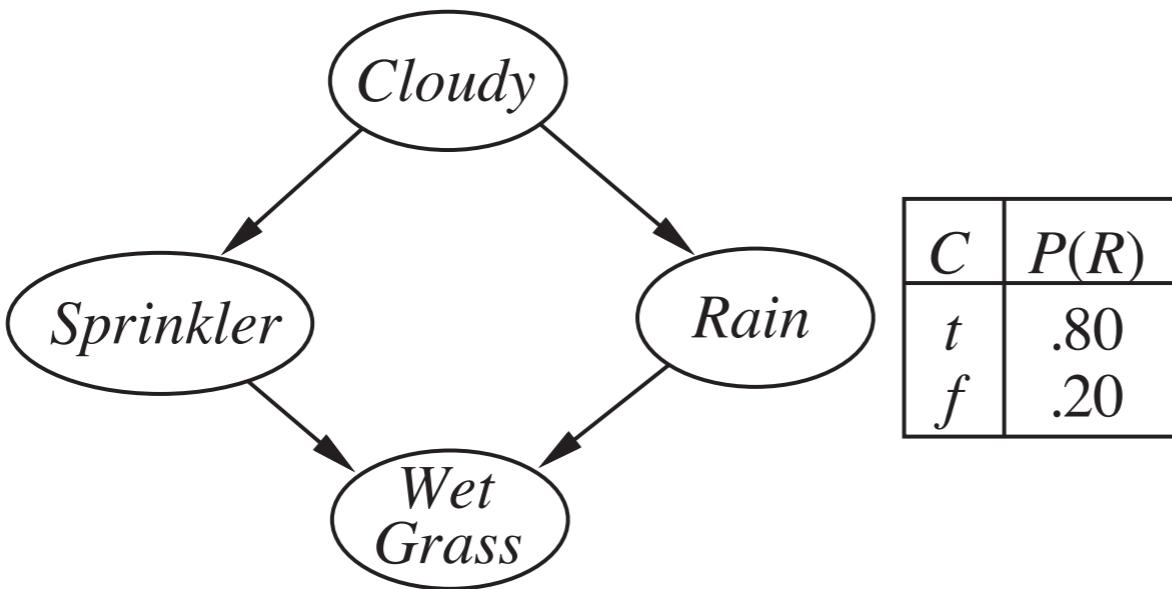
C	P(S)
t	.10
f	.50

C	P(R)
t	.80
f	.20

S	R	P(W)
t	t	.99
t	f	.90
f	t	.90
f	f	.00

$$P(C) = .5$$

C	P(S)
t	.10
f	.50



C	P(R)
t	.80
f	.20

S	R	P(W)
t	t	.99
t	f	.90
f	t	.90
f	f	.00

Cloudy

Sprinkler

Rain

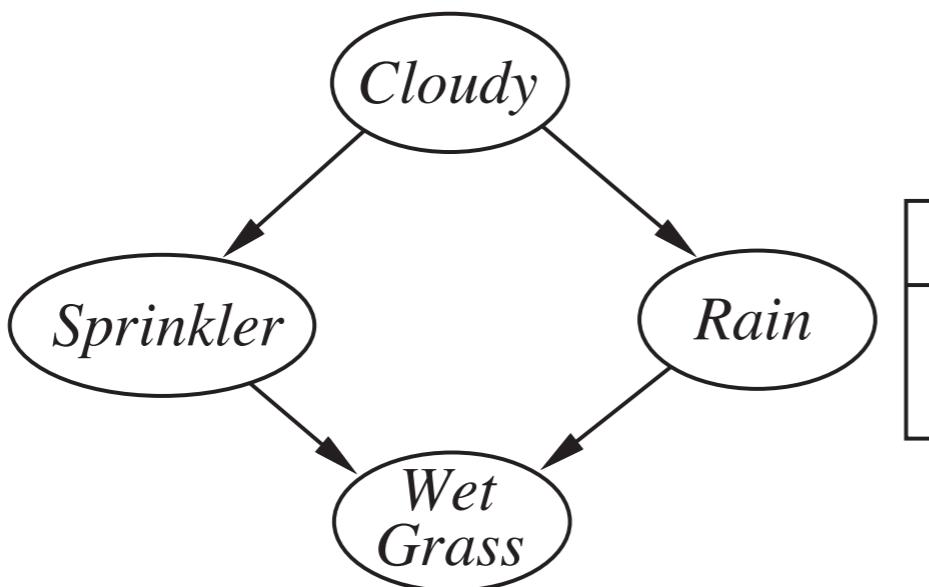
WetGrass

true

$$\mathbf{P}(Cloudy) = \langle 0.5, 0.5 \rangle$$

C	$P(S)$
t	.10
f	.50

$$P(C) = .5$$



C	$P(R)$
t	.80
f	.20

S	R	$P(W)$
t	t	.99
t	f	.90
f	t	.90
f	f	.00

Cloudy

Sprinkler

Rain

WetGrass

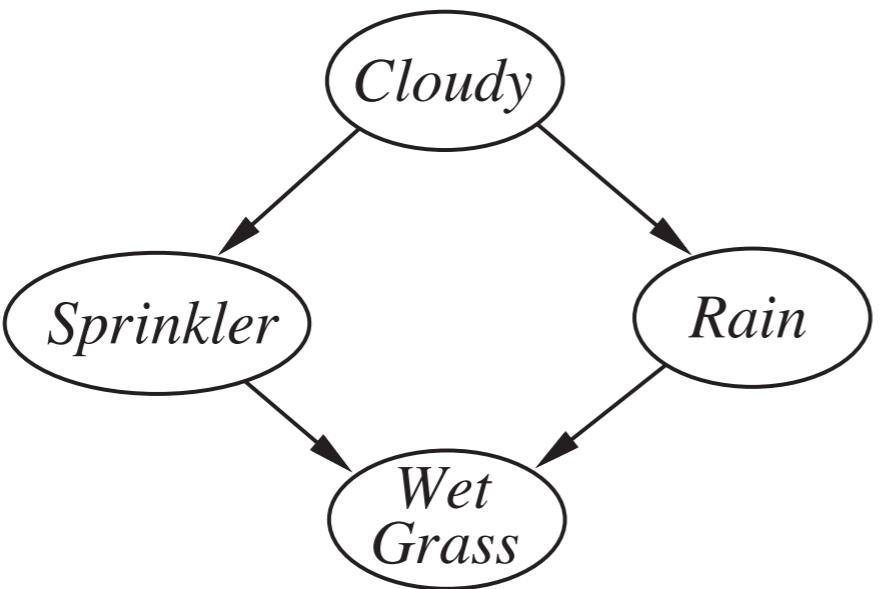
true

false

$$\mathbf{P}(\textit{Sprinkler} \mid \textit{Cloudy} = \textit{true}) = \langle 0.1, 0.9 \rangle$$

C	$P(S)$
t	.10
f	.50

$$P(C) = .5$$



C	$P(R)$
t	.80
f	.20

S	R	$P(W)$
t	t	.99
t	f	.90
f	t	.90
f	f	.00

Cloudy

Sprinkler

Rain

WetGrass

true

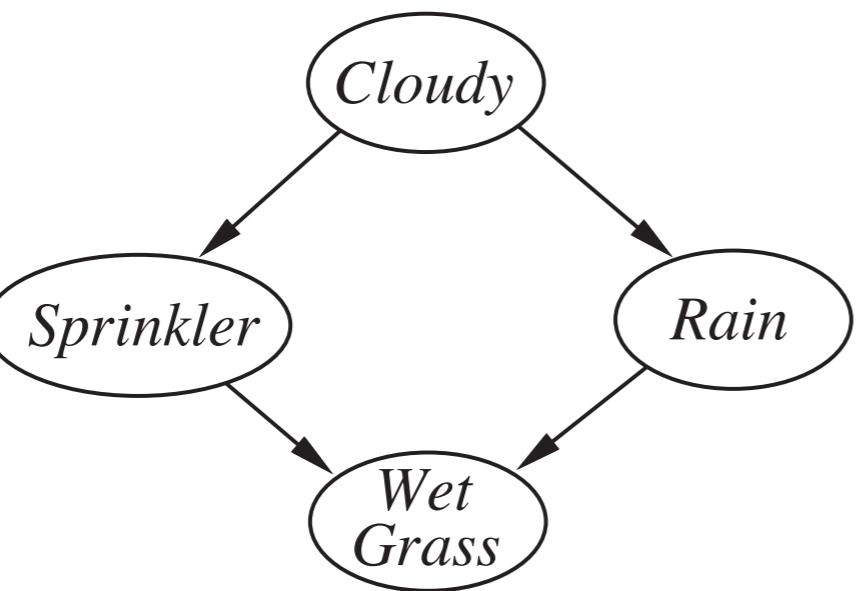
false

true

$$\mathbf{P}(Rain \mid Cloudy = true) = \langle 0.8, 0.2 \rangle$$

C	$P(S)$
t	.10
f	.50

$$P(C) = .5$$



C	$P(R)$
t	.80
f	.20

S	R	$P(W)$
t	t	.99
t	f	.90
f	t	.90
f	f	.00

Cloudy

Sprinkler

Rain

WetGrass

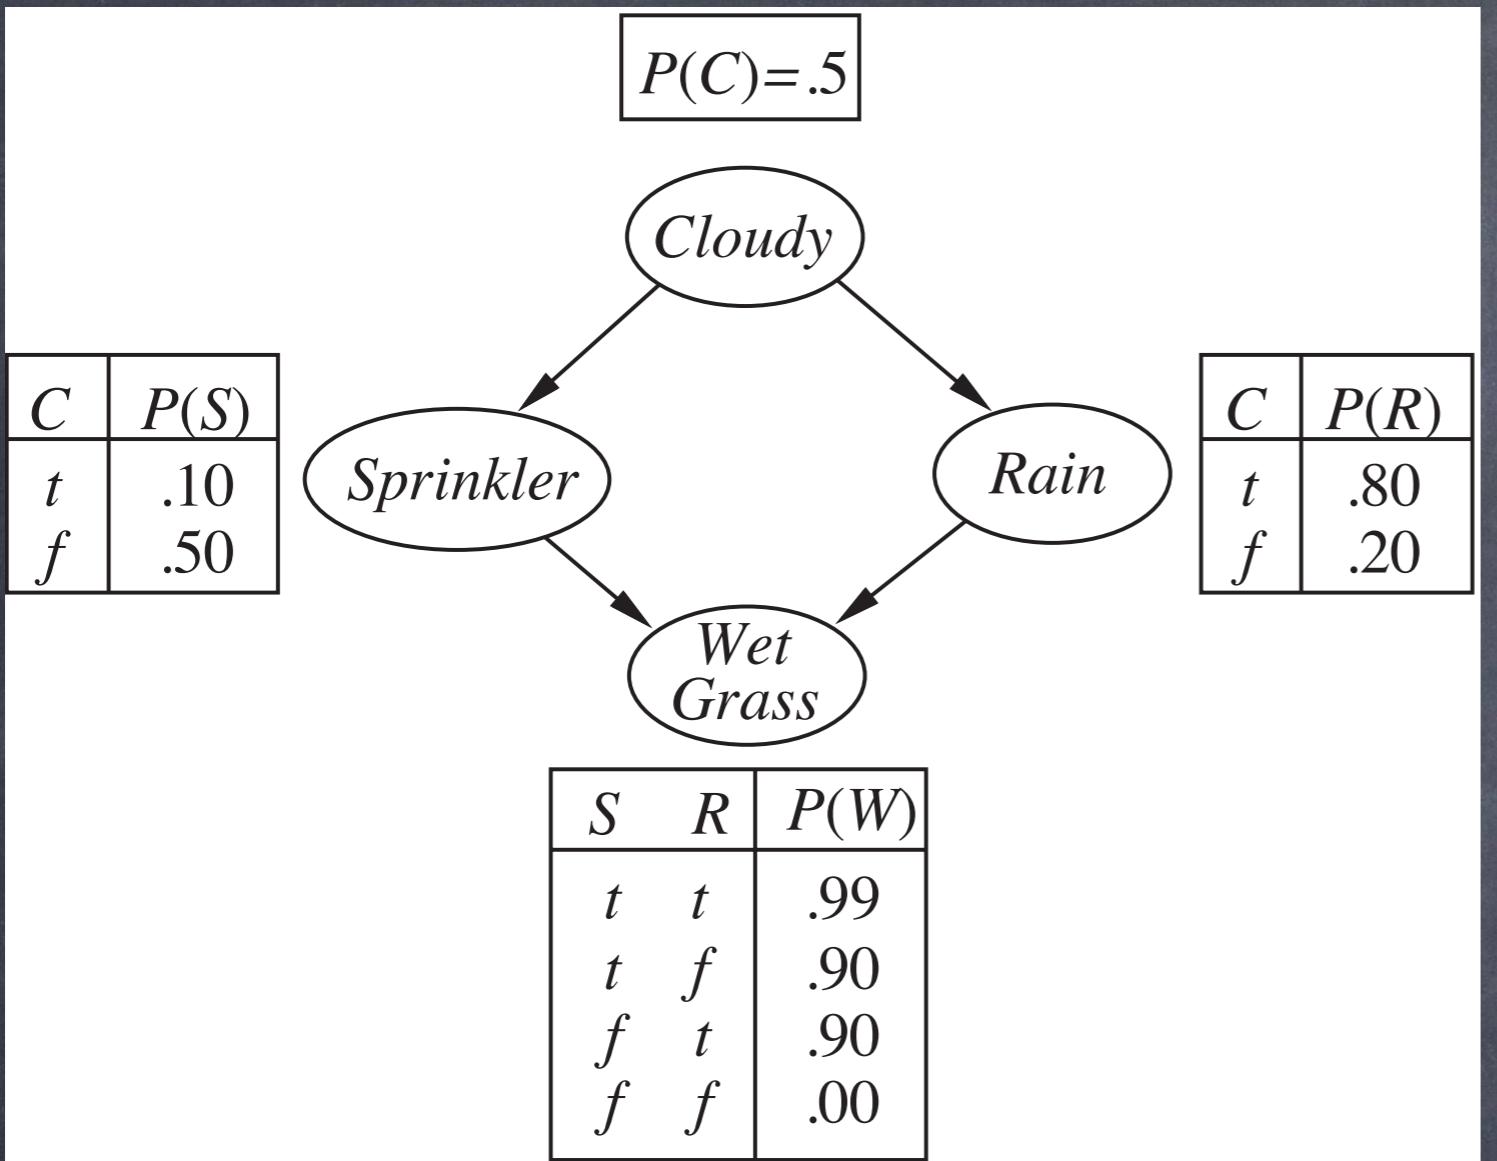
true

false

true

true

$$\mathbf{P}(WetGrass \mid Sprinkler = false, Rain = true) = \langle 0.9, 0.1 \rangle$$



Cloudy	true
Sprinkler	false
Rain	true
WetGrass	true

$\langle \text{Cloudy}=\text{true}, \text{Sprinkler}=\text{false}, \text{Rain}=\text{true}, \text{WetGrass}=\text{true} \rangle$

Guaranteed to be a consistent estimate
(becomes exact in the large-sample limit)

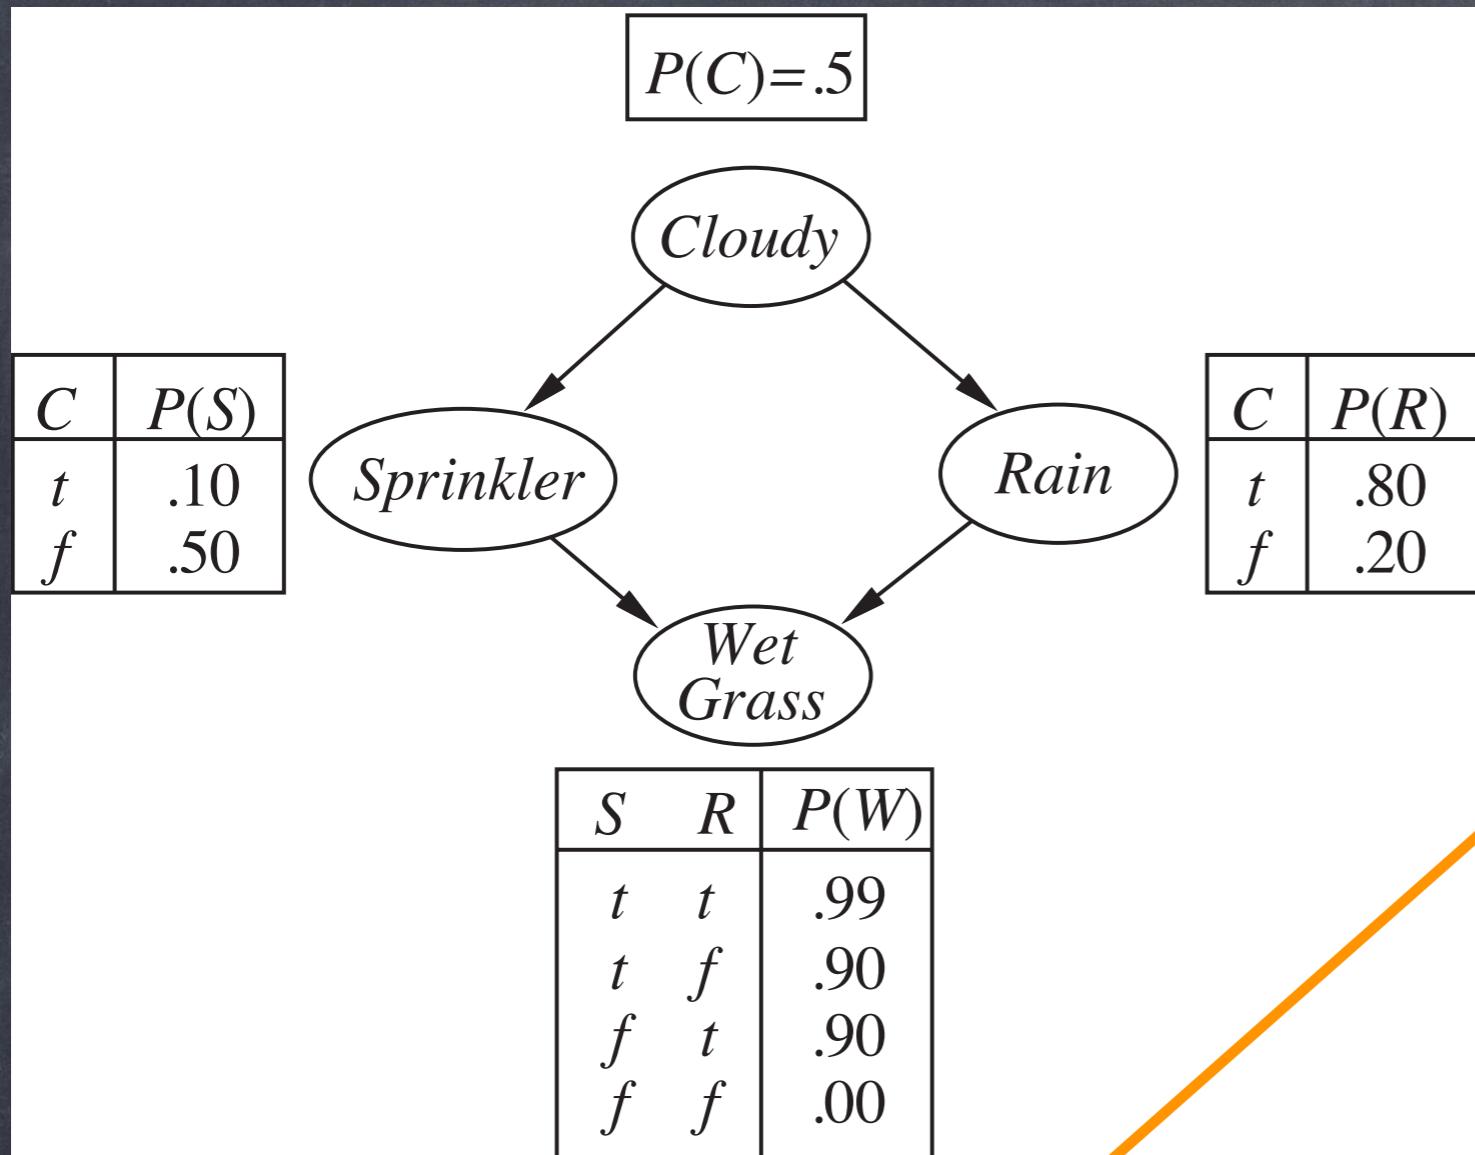
Sampling

- Generate assignments of values to the random variables
- That are consistent with the full joint distribution encoded in the network
 - In the sense that in the limit, the frequency of occurrence of any event (possible world) is equal to its probability

Approximate Inference

- Query variable X
- Evidence variables E_1, \dots, E_m
 - Observed values: $e = \langle e_1, \dots, e_m \rangle$
- Non-evidence, non-query (“hidden”) variables: Y
- Approximate (estimate): $P(X | e)$

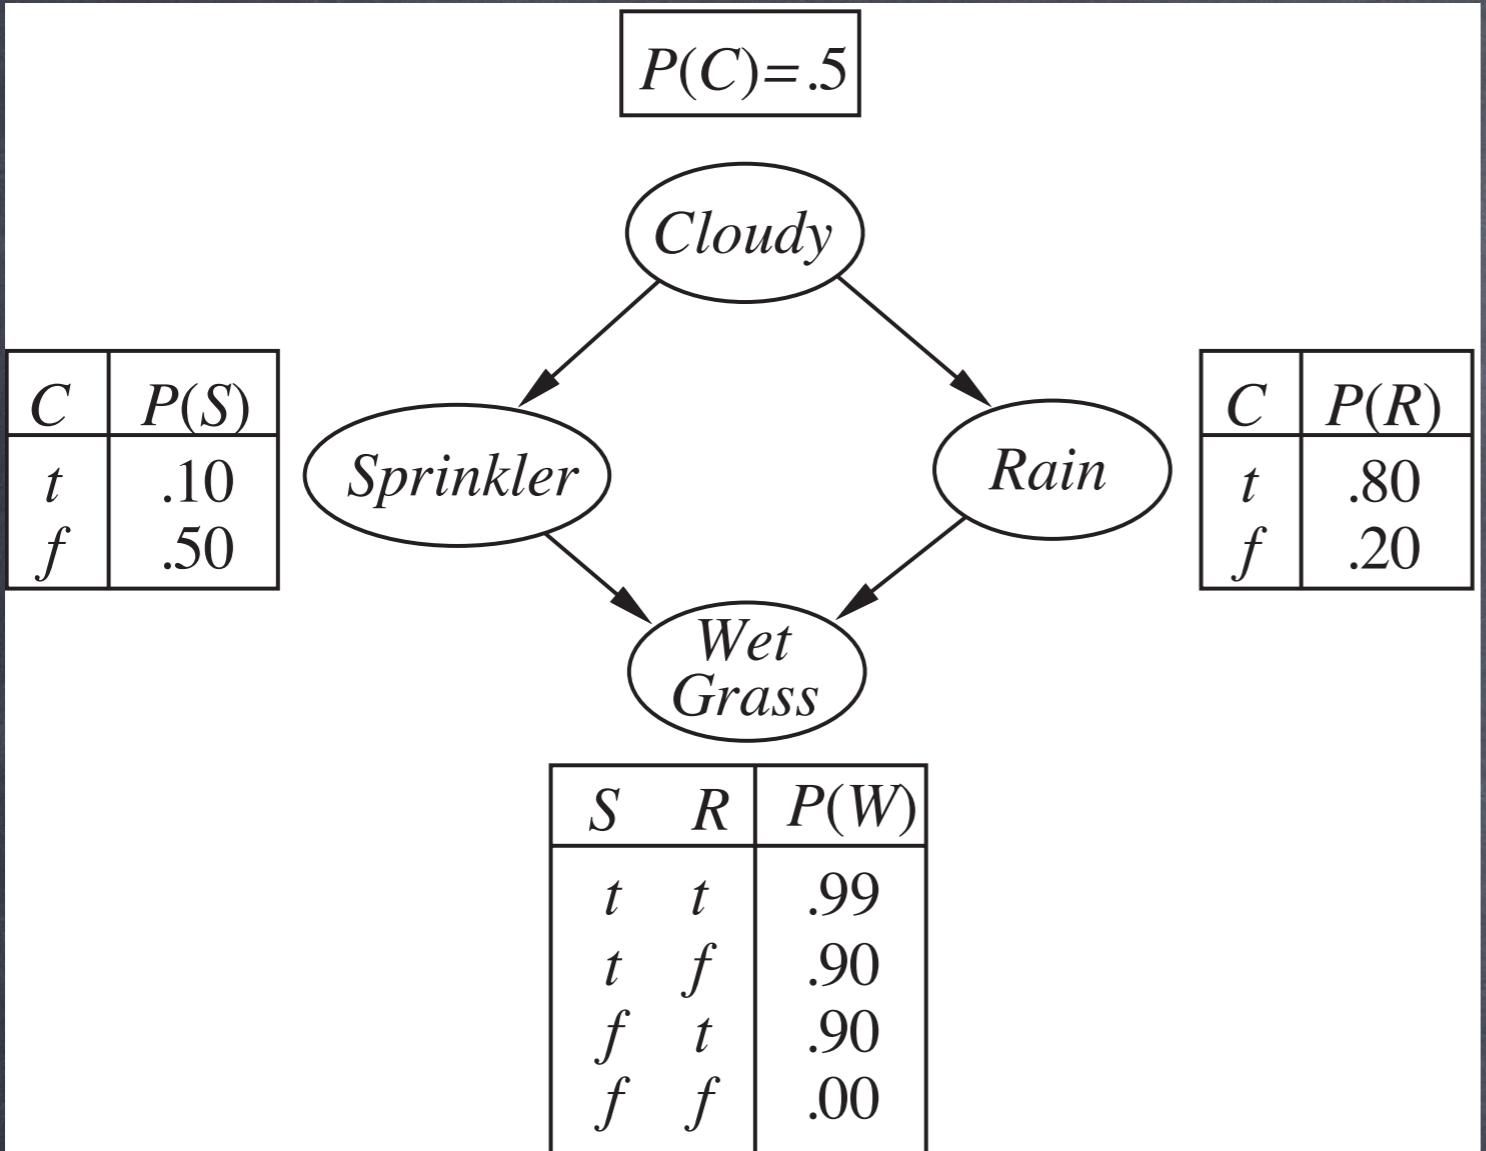
$$\mathbf{P}(Rain \mid Sprinkler = true)$$



\langle Cloudy=true, Sprinkler=false, Rain=true, WetGrass=true \rangle

Rejection Sampling

- Generate sample from the prior distribution specified by the network
- Reject sample if inconsistent with the evidence
- Use remaining samples to estimate probability of event

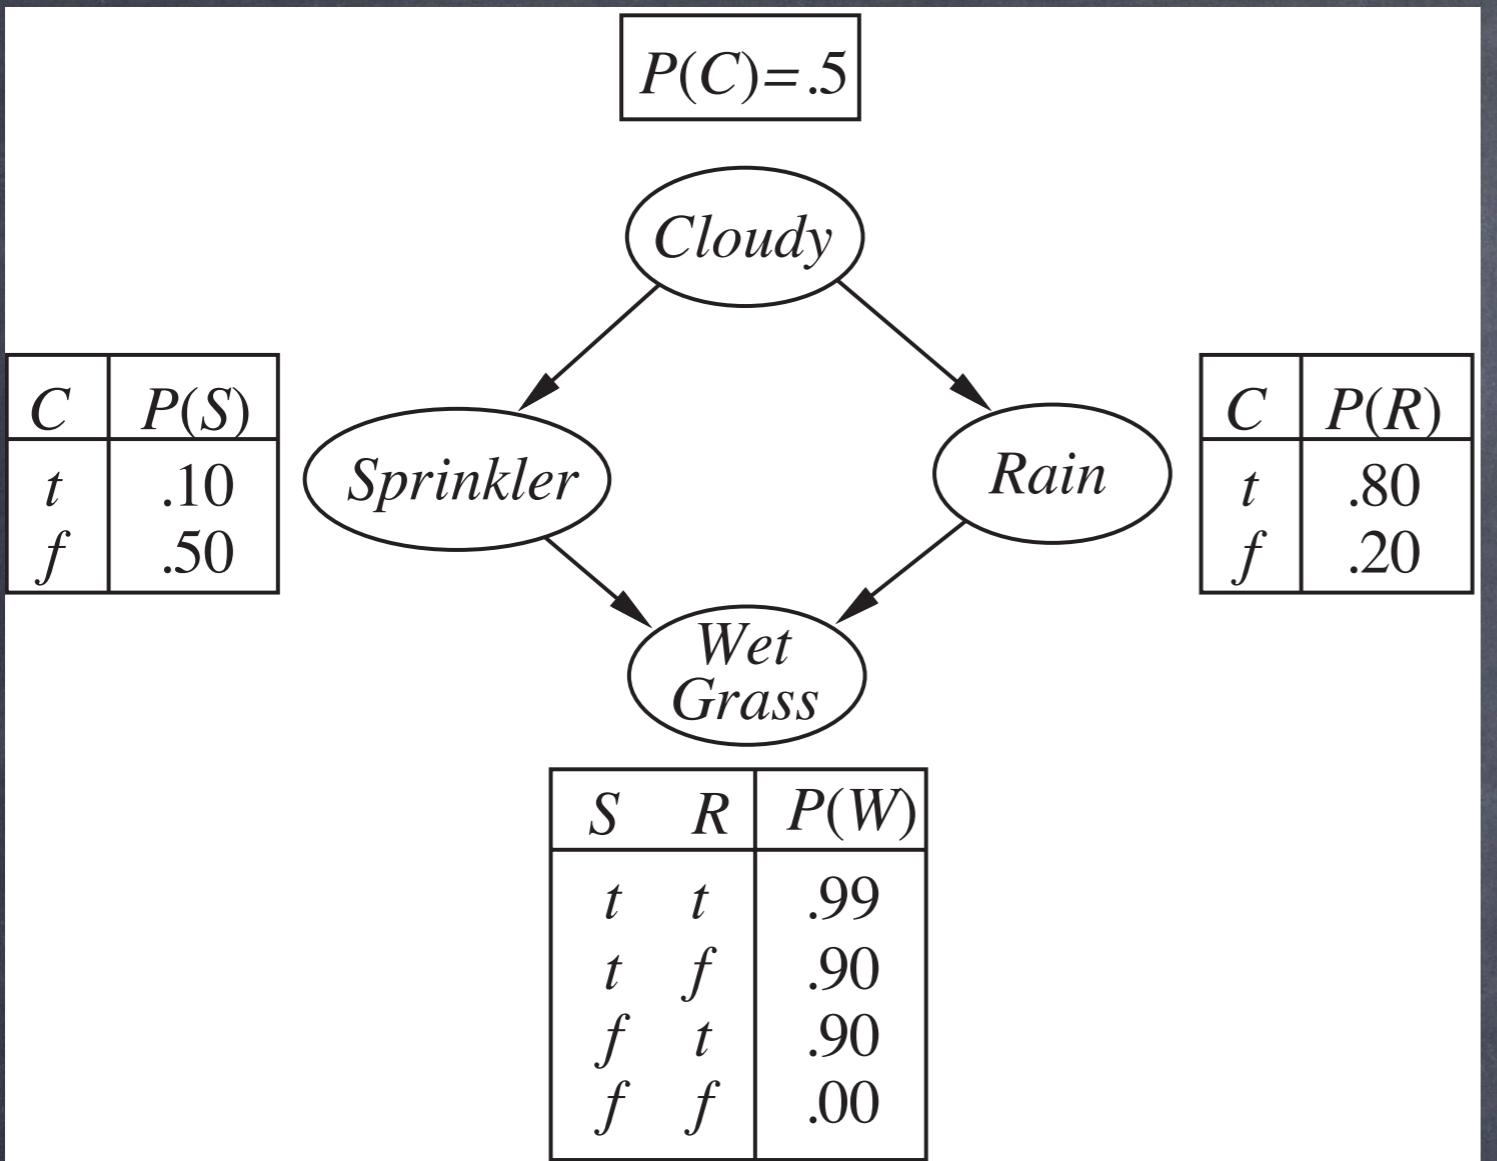


$P(Rain \mid Sprinkler = \text{true})$

100 samples

Sprinkler=false: 73

Sprinkler=true: 27



$\mathbf{P}(Rain \mid Sprinkler = true)$

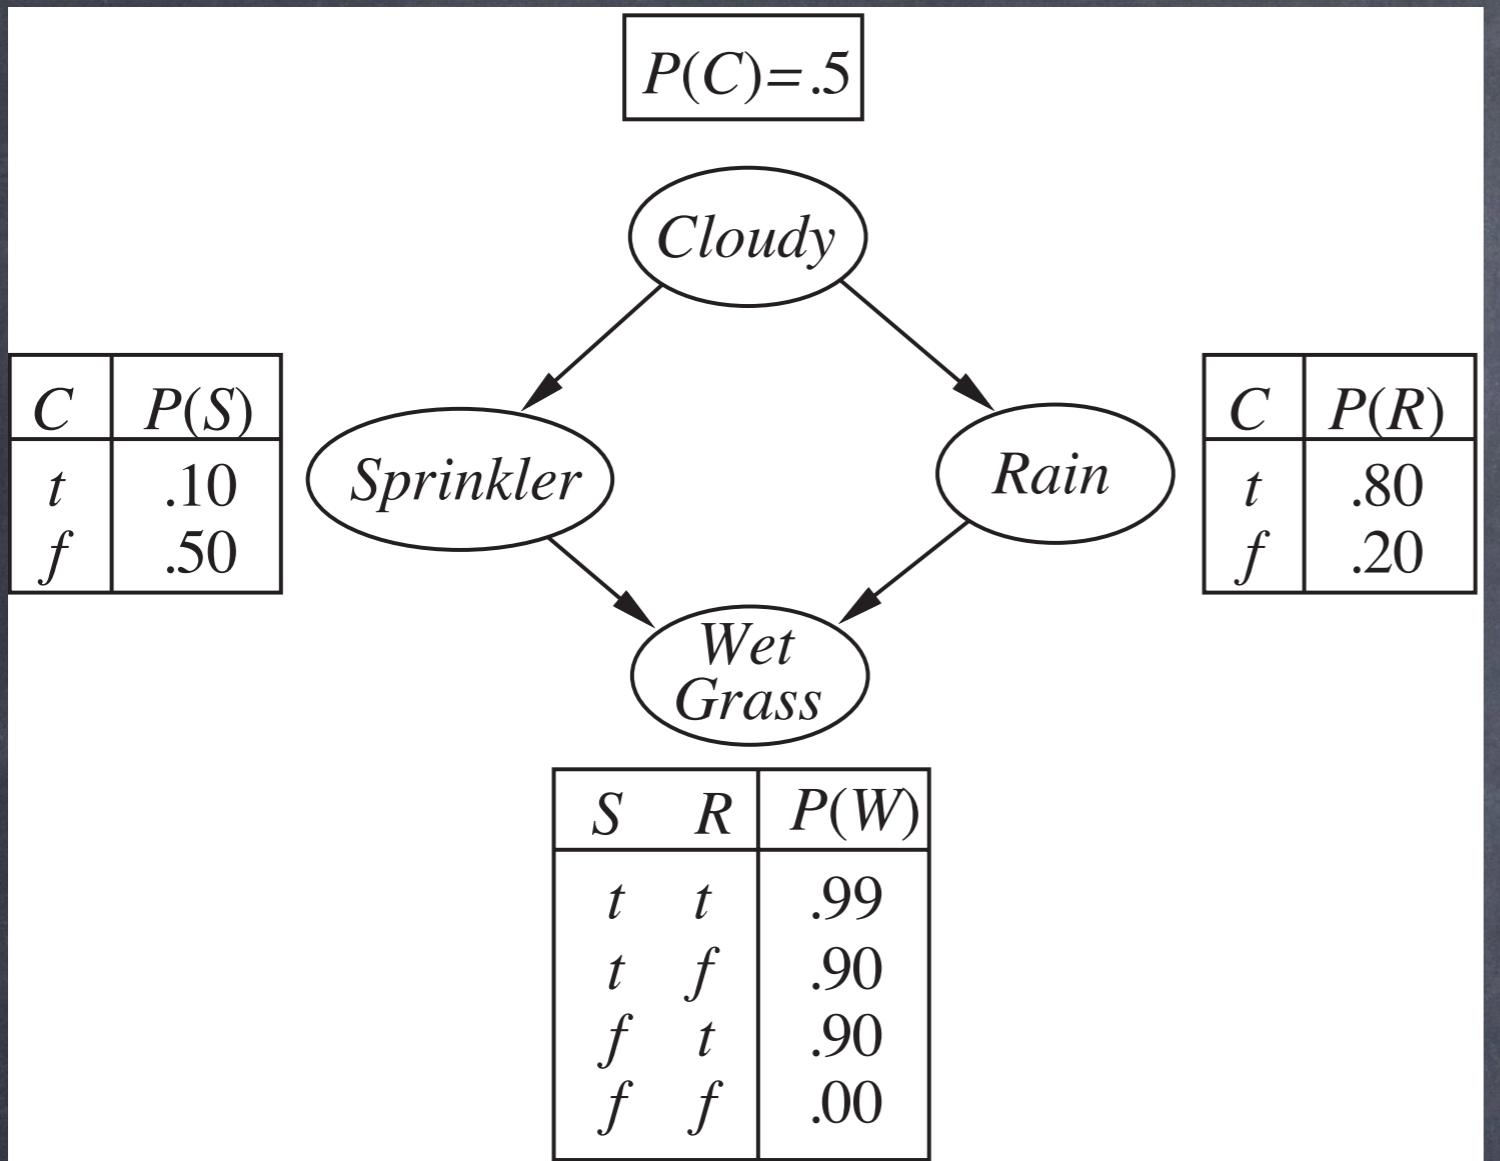
100 samples

Sprinkler=false: 73

Sprinkler=true: 27

Rain=true: 8

Rain=false: 19



$$\mathbf{P}(Rain \mid Sprinkler = true)$$

100 samples

Sprinkler=false: 73

Sprinkler=true: 27

Rain=true: 8

Rain=false: 19

$$\mathbf{P}(Rain \mid Sprinkler = true) \approx \alpha \left\langle \frac{8}{27}, \frac{19}{27} \right\rangle = \langle 0.296, 0.704 \rangle$$

Rejection Sampling

- Generate sample from the prior distribution specified by the network
- Reject sample if inconsistent with the evidence
- Use remaining samples to estimate probability of event

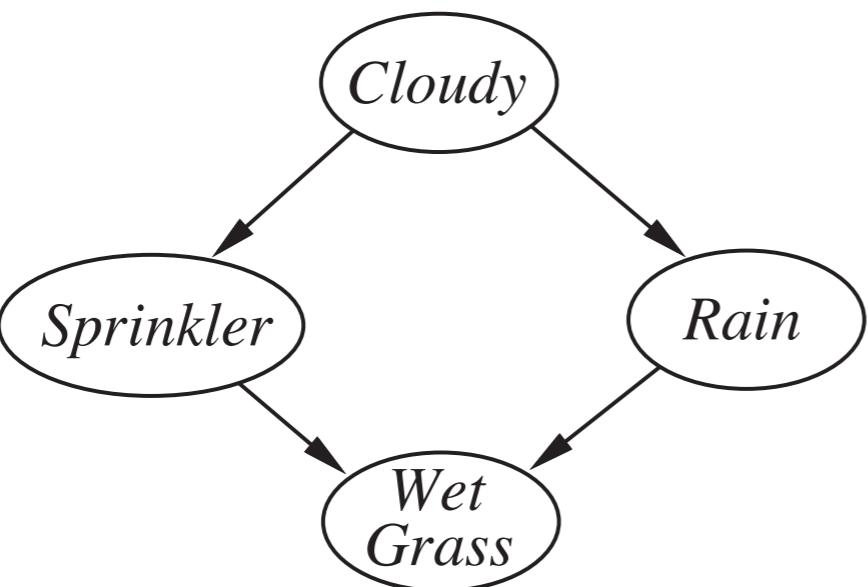
Rejection Sampling

- Generate sample from the prior distribution specified by the network
- Reject sample if inconsistent with the evidence
- Use remaining samples to estimate probability of event

Fraction of samples consistent with the evidence drops exponentially with number of evidence variables

Cloudy
Sprinkler
Rain
WetGrass

$$P(C) = .5$$

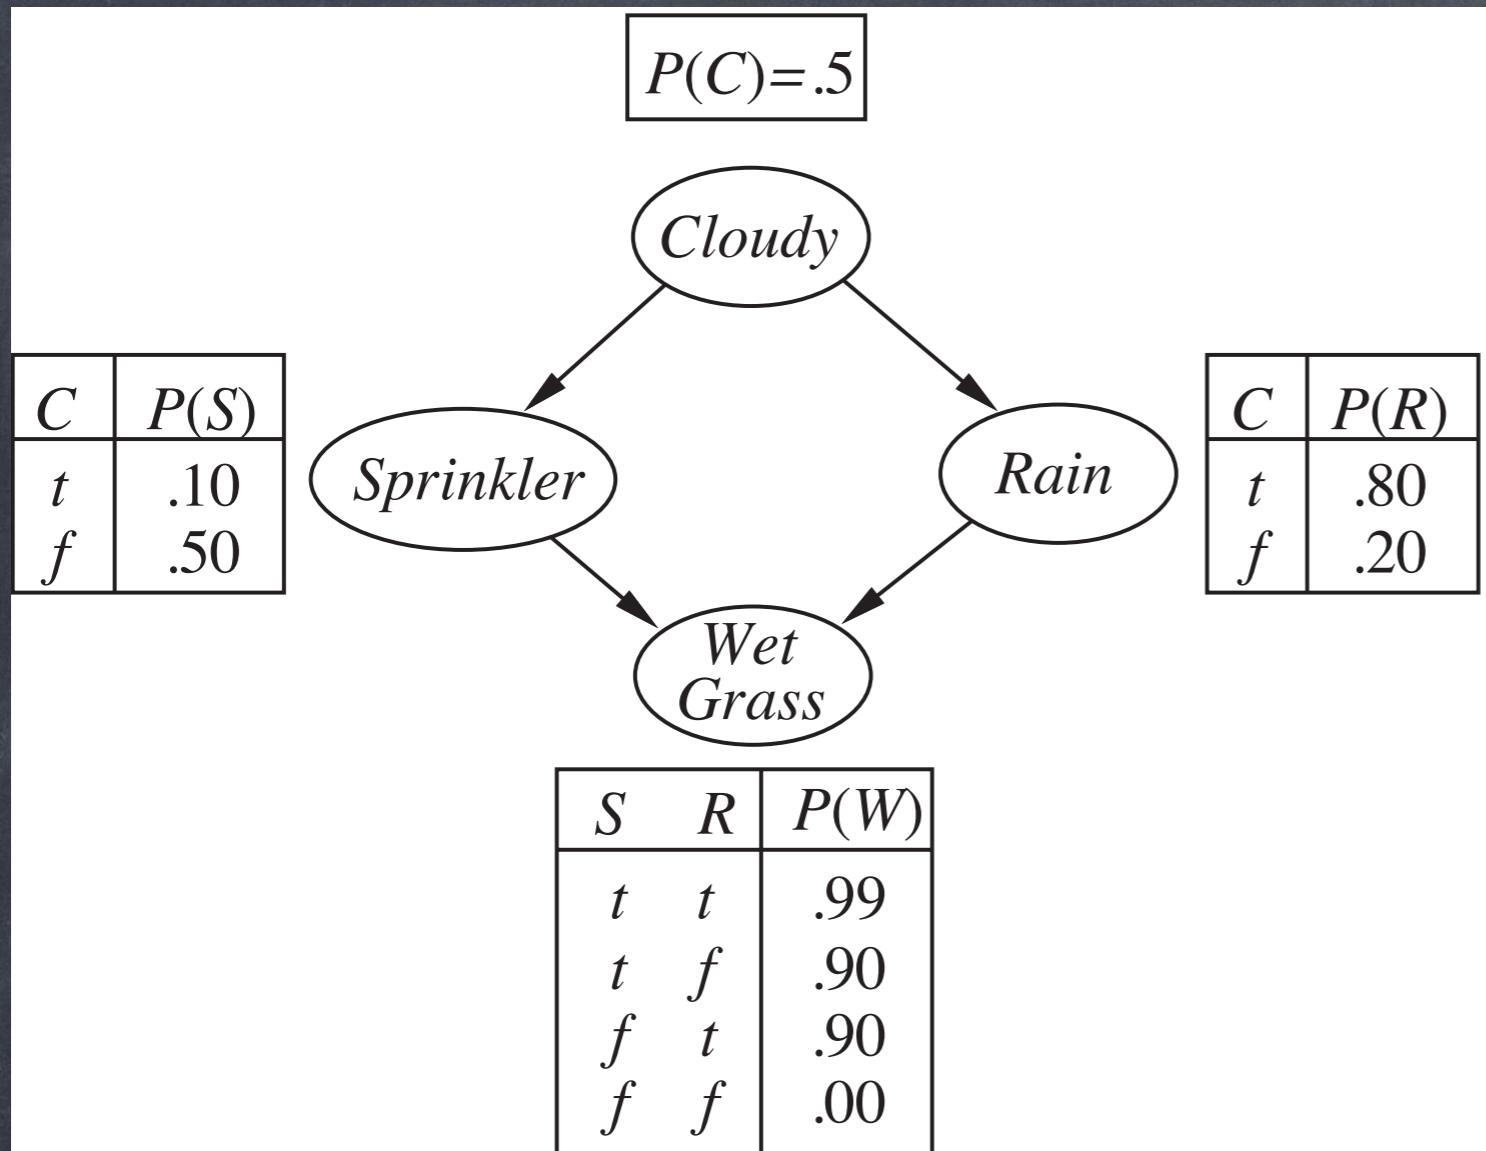


C	P(S)
t	.10
f	.50

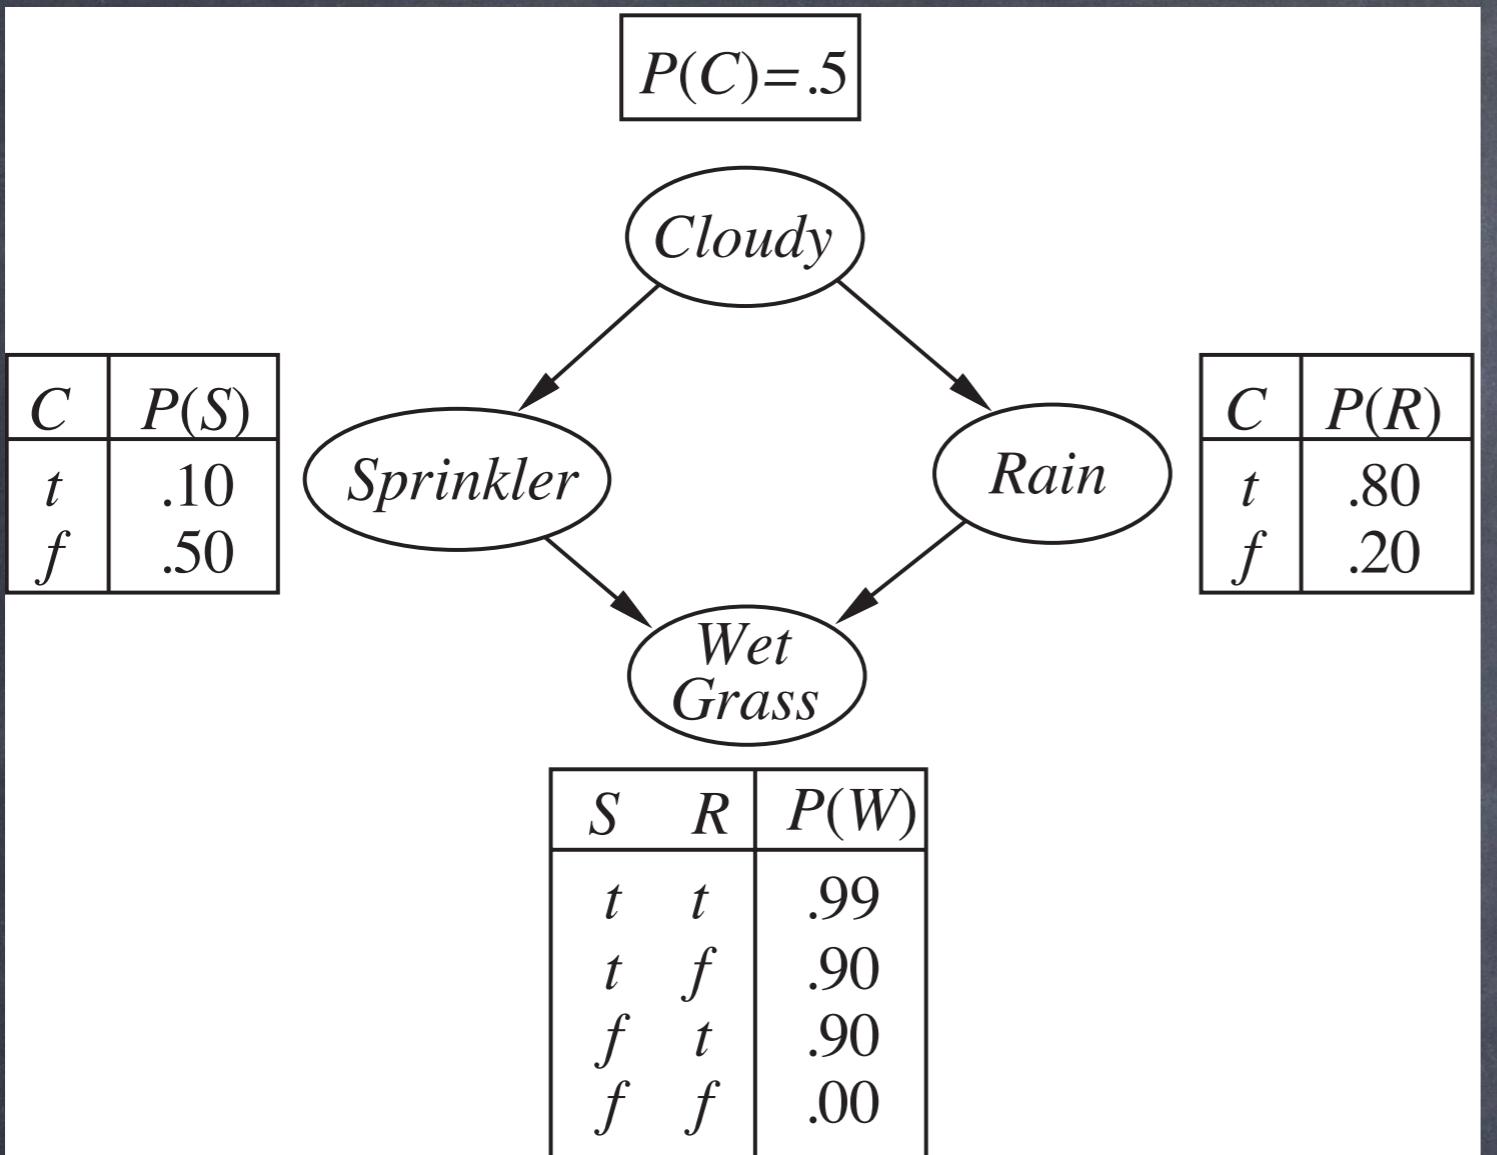
C	P(R)
t	.80
f	.20

S	R	P(W)
t	t	.99
t	f	.90
f	t	.90
f	f	.00

Cloudy
Sprinkler
Rain
WetGrass



$$P(Rain \mid Cloudy = \text{true}, WetGrass = \text{true})$$



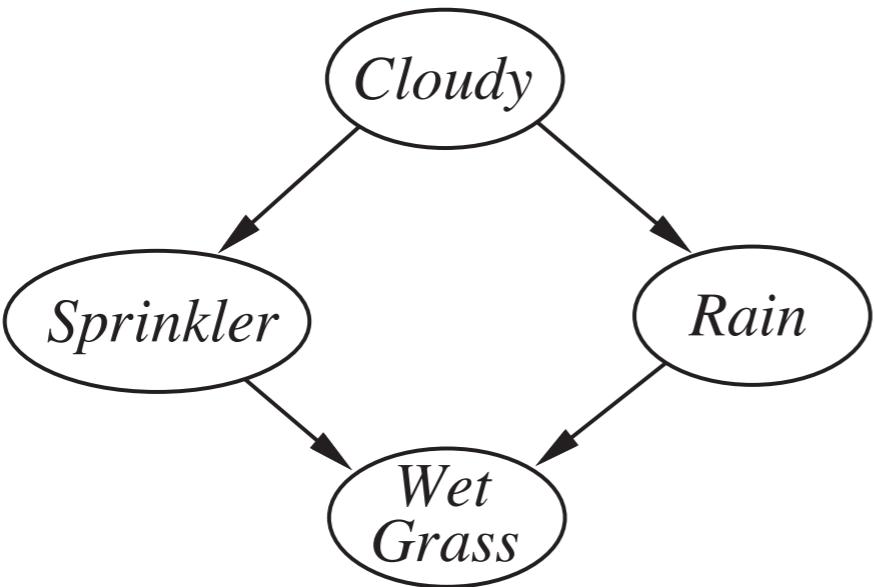
Cloudy
Sprinkler
Rain
WetGrass

$$w = 1.0$$

$$\mathbf{P}(Rain \mid Cloudy = true, WetGrass = true)$$

C	$P(S)$
t	.10
f	.50

$$P(C) = .5$$



C	$P(R)$
t	.80
f	.20

S	R	$P(W)$
t	t	.99
t	f	.90
f	t	.90
f	f	.00

Cloudy

Sprinkler

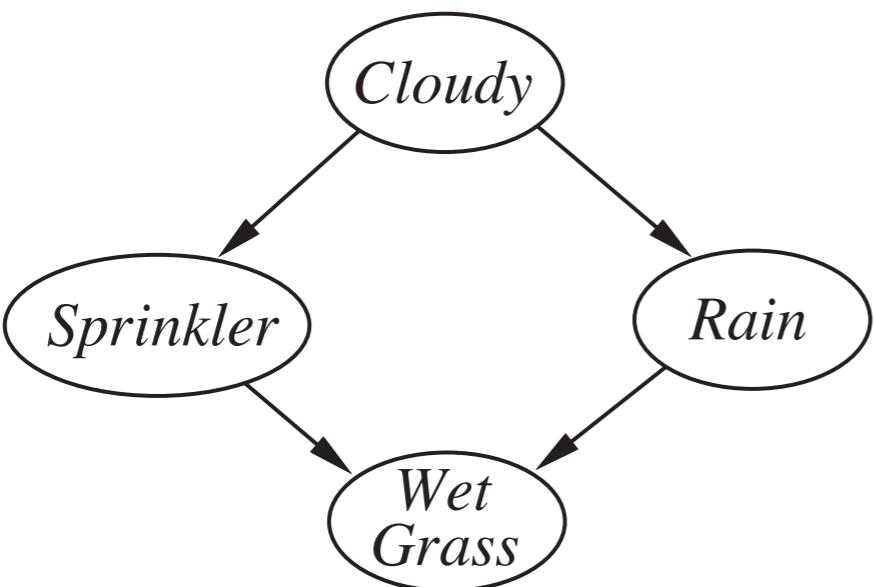
Rain

WetGrass

W = 1.0

$$\mathbf{P}(Rain \mid Cloudy = true, WetGrass = true)$$

$$P(C) = .5$$



C	$P(S)$
t	.10
f	.50

C	$P(R)$
t	.80
f	.20

S	R	$P(W)$
t	t	.99
t	f	.90
f	t	.90
f	f	.00

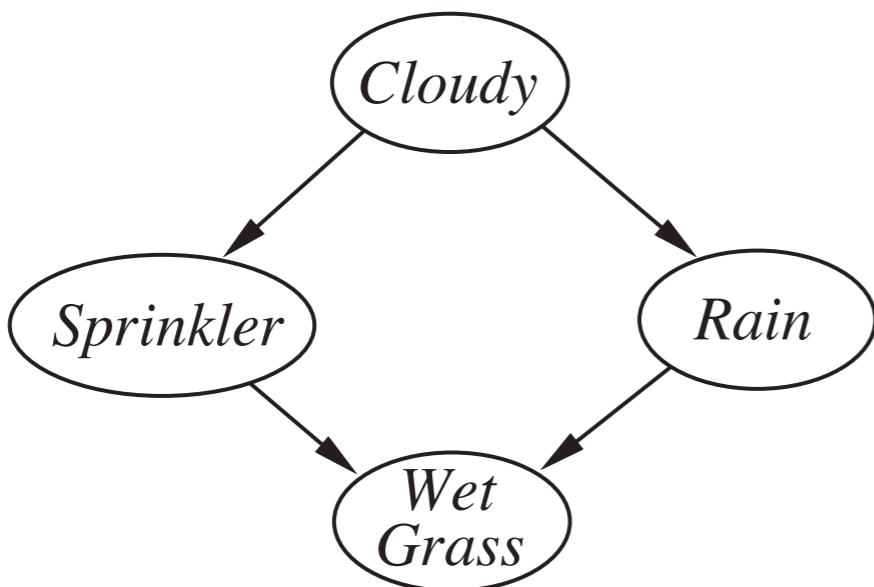
Cloudy true
Sprinkler
Rain
WetGrass

w = 0.5

$$\mathbf{P}(Rain \mid Cloudy = true, WetGrass = true)$$

C	$P(S)$
t	.10
f	.50

$$P(C)=.5$$



C	$P(R)$
t	.80
f	.20

S	R	$P(W)$
t	t	.99
t	f	.90
f	t	.90
f	f	.00

Cloudy

Sprinkler

Rain

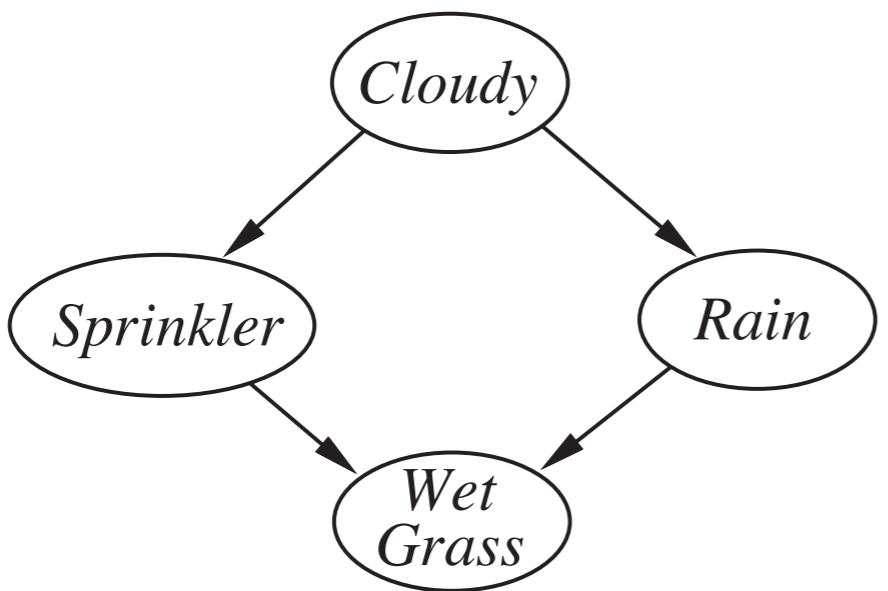
WetGrass

W = 0.5

$$\mathbf{P}(Rain \mid Cloudy = true, WetGrass = true)$$

C	$P(S)$
t	.10
f	.50

$$P(C)=.5$$



C	$P(R)$
t	.80
f	.20

S	R	$P(W)$
t	t	.99
t	f	.90
f	t	.90
f	f	.00

Cloudy

Sprinkler

Rain

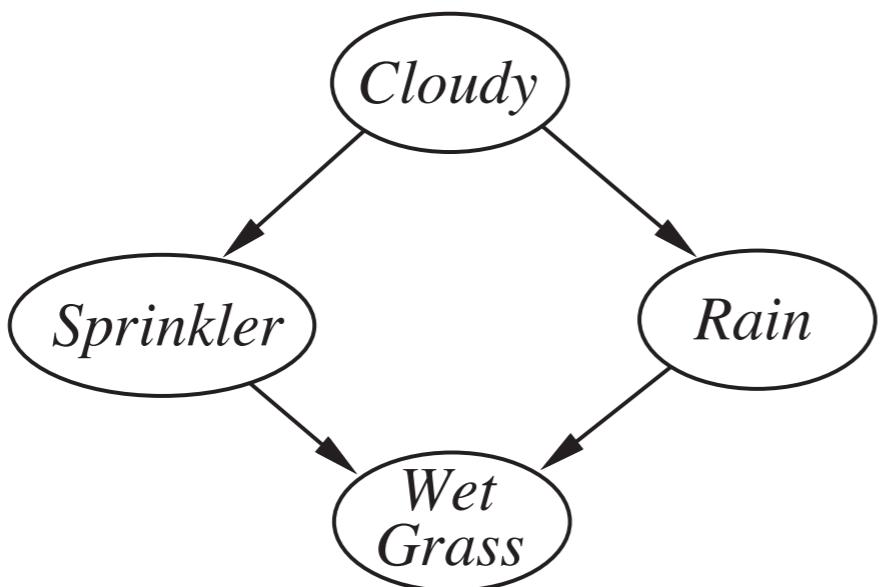
WetGrass

W = 0.5

$$\mathbf{P}(Rain \mid Cloudy = true, WetGrass = true)$$

C	$P(S)$
t	.10
f	.50

$$P(C)=.5$$



C	$P(R)$
t	.80
f	.20

S	R	$P(W)$
t	t	.99
t	f	.90
f	t	.90
f	f	.00

Cloudy

Sprinkler

Rain

WetGrass

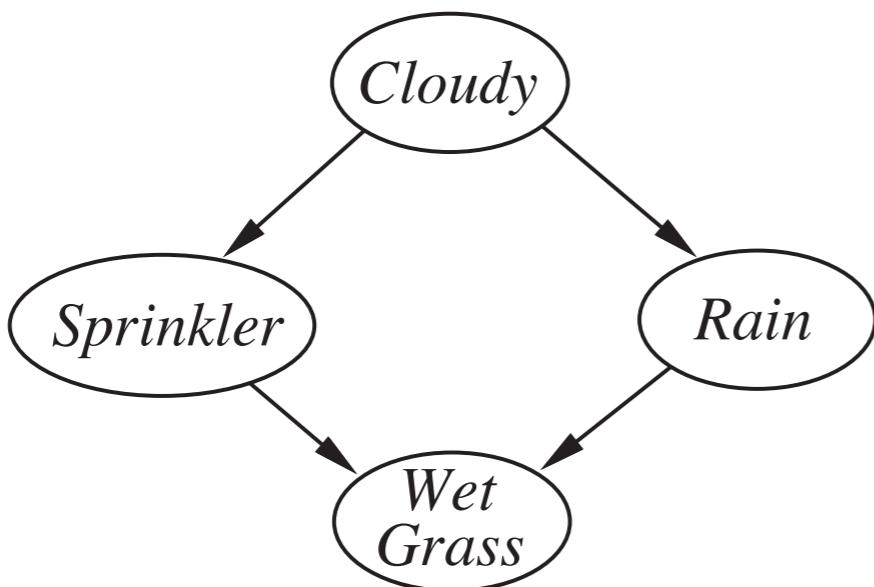
W = 0.5

true
false
true
true

$$\mathbf{P}(Rain \mid Cloudy = true, WetGrass = true)$$

C	$P(S)$
t	.10
f	.50

$$P(C) = .5$$



C	$P(R)$
t	.80
f	.20

S	R	$P(W)$
t	t	.99
t	f	.90
f	t	.90
f	f	.00

Cloudy

Sprinkler

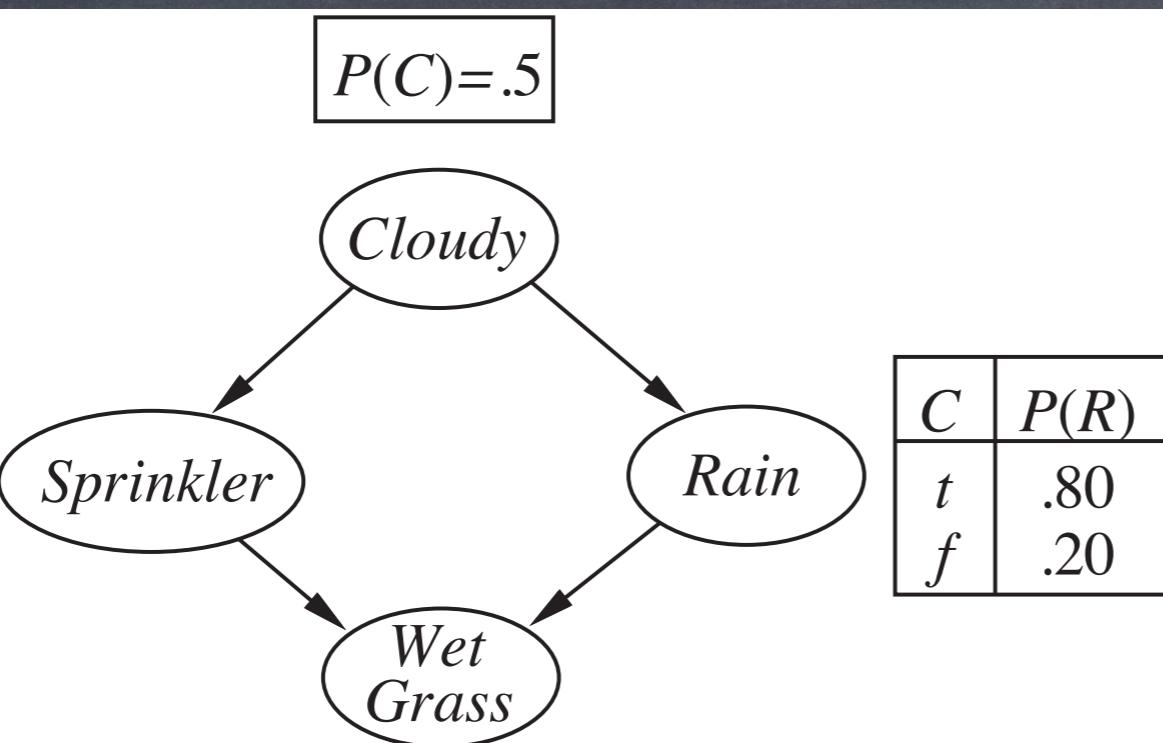
Rain

WetGrass

W = 0.45

$$\text{P}(Rain \mid Cloudy = true, WetGrass = true)$$

C	$P(S)$
t	.10
f	.50



S	R	$P(W)$
t	t	.99
t	f	.90
f	t	.90
f	f	.00

Cloudy	true
Sprinkler	false
Rain	true
WetGrass	true

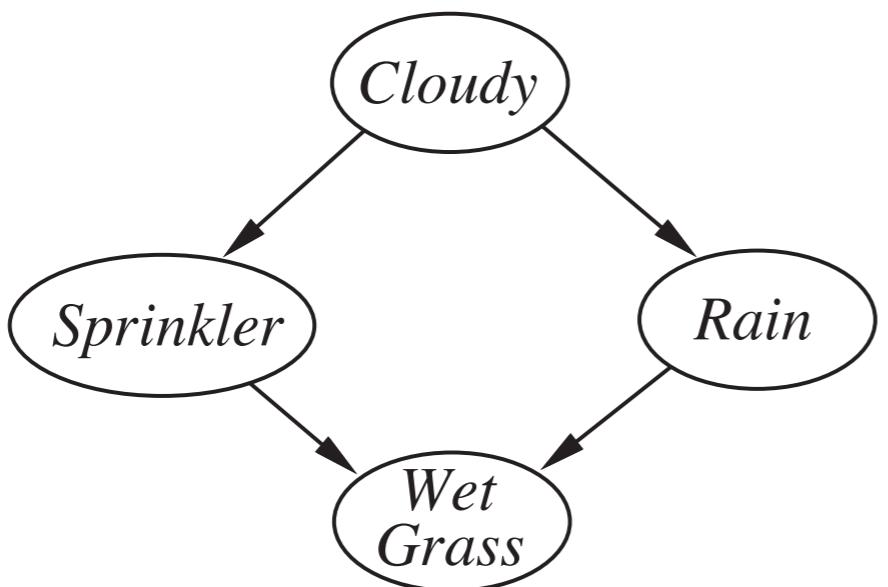
W = 0.45

$$\mathbf{P}(Rain \mid Cloudy = true, WetGrass = true)$$

N Rain=true +=

C	$P(S)$
t	.10
f	.50

$$P(C)=.5$$



C	$P(R)$
t	.80
f	.20

S	R	$P(W)$
t	t	.99
t	f	.90
f	t	.90
f	f	.00

Cloudy	true
Sprinkler	false
Rain	true
WetGrass	true

W = 0.45

$$\mathbf{P}(Rain \mid Cloudy = true, WetGrass = true)$$

$N_{Rain=true} += 0.45$

Likelihood Weighting

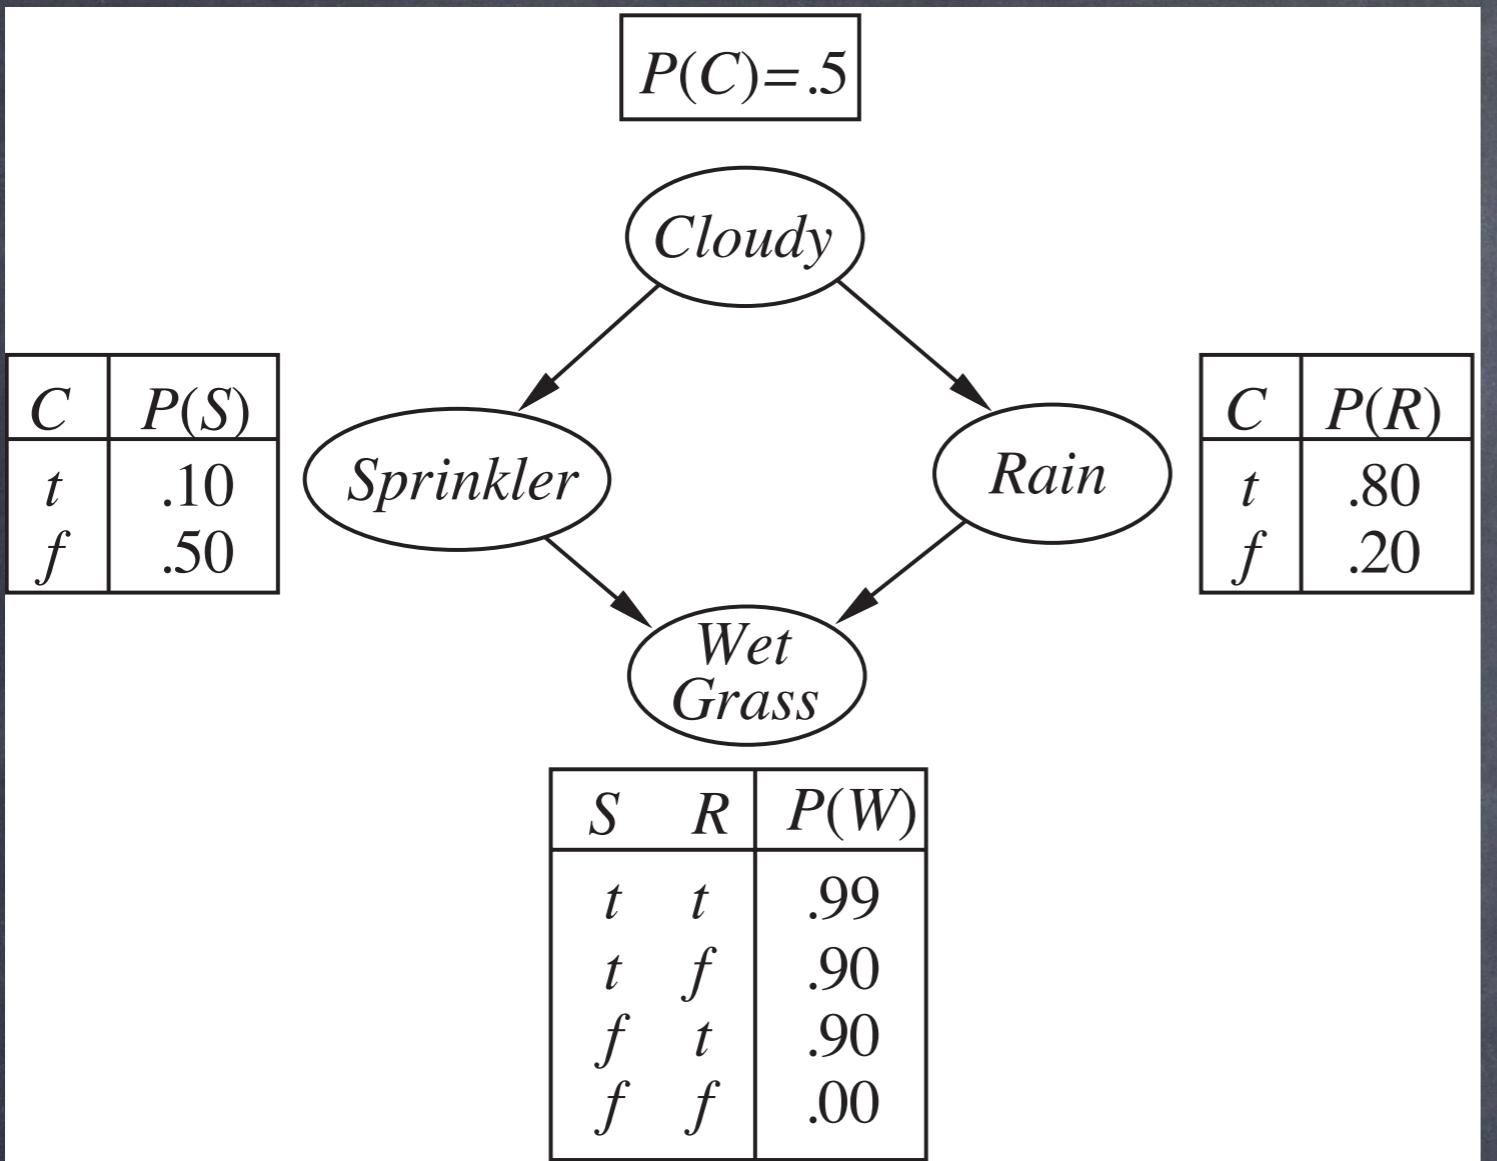
- Generate only samples consistent with the evidence
 - Generate sample using topological order
 - Evidence variable: Value fixed, update weight
 - Non-evidence variable: Sample value from network
- Compute weighted sum of samples for estimate

Likelihood Weighting

- Pros:
 - Doesn't reject any samples
- Cons:
 - More evidence \Rightarrow lower weight
 - Affected by order of evidence vars in topsort (later = worse)

Approximate Inference in Bayesian Networks

- Rejection Sampling
- Likelihood Weighting



Cloudy	true
Sprinkler	false
Rain	true
WetGrass	true

$$P(Rain \mid Sprinkler = \text{true}, WetGrass = \text{true})$$

- To approximate: $P(X | e)$
- Generate a sequence of states
 - Values of evidence variables are fixed
 - Values of other variables appear in the right proportion given the distribution encoded by the network

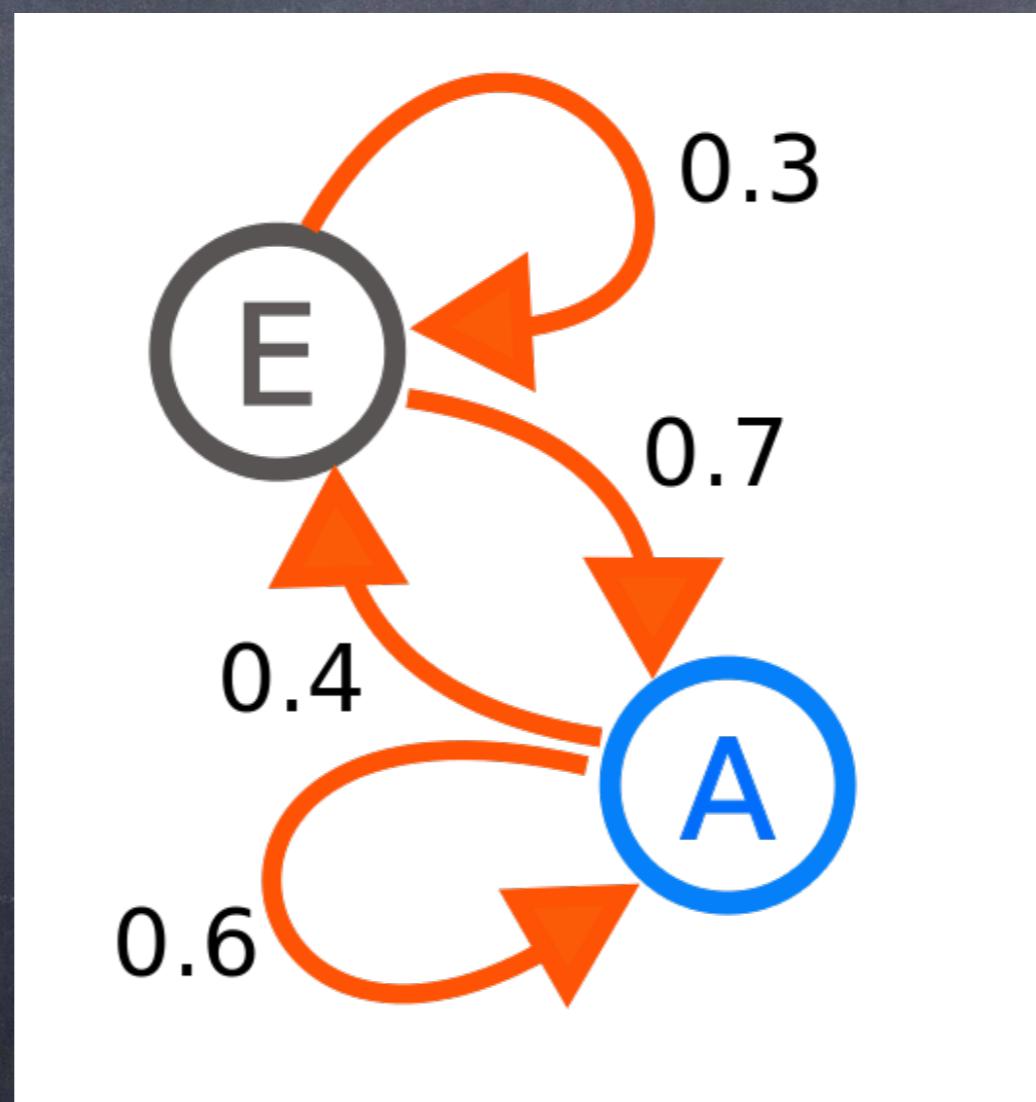
Markov Chain Monte Carlo Simulation

- To approximate: $P(X | e)$
- Generate a sequence of states
 - Values of evidence variables are fixed
 - Values of other variables appear in the right proportion given the distribution encoded by the network

Monte Carlo

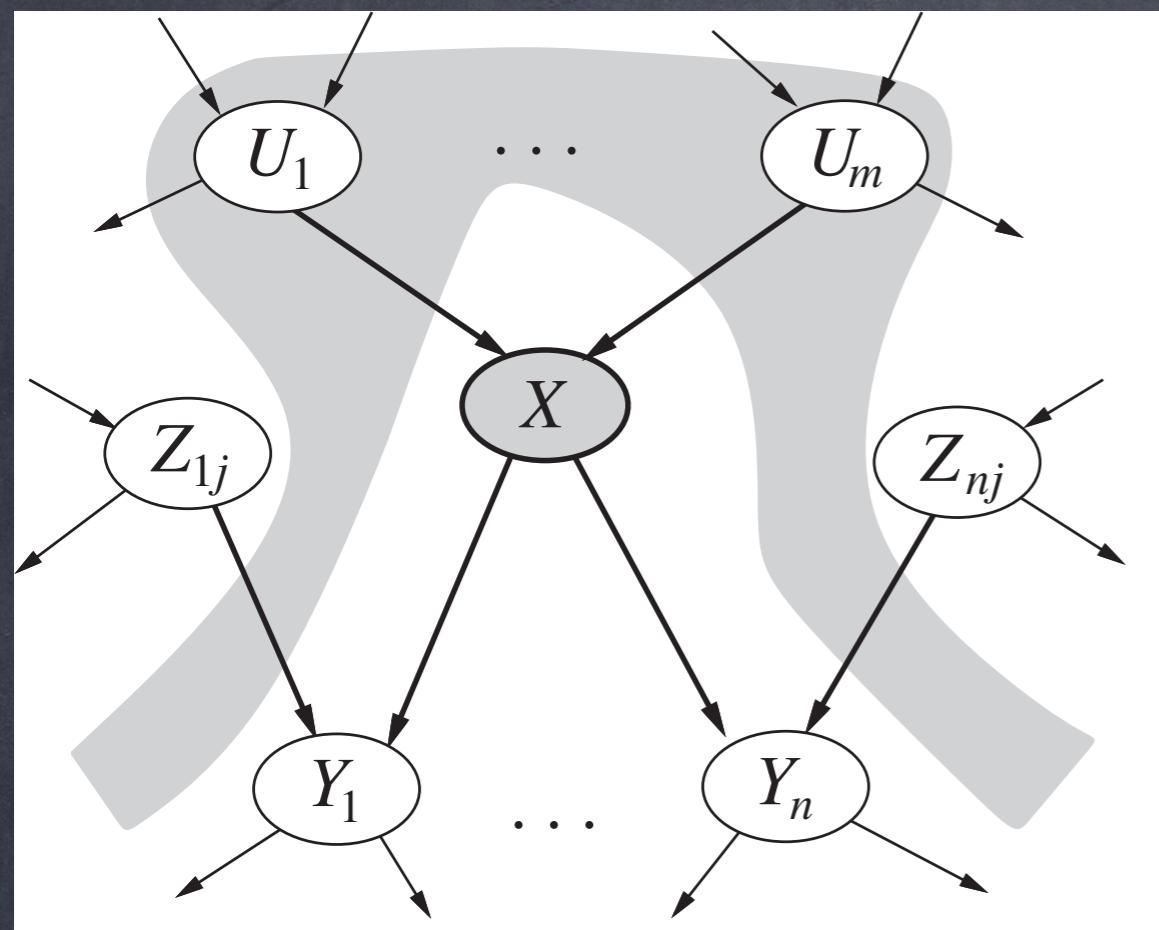


Markov Chain



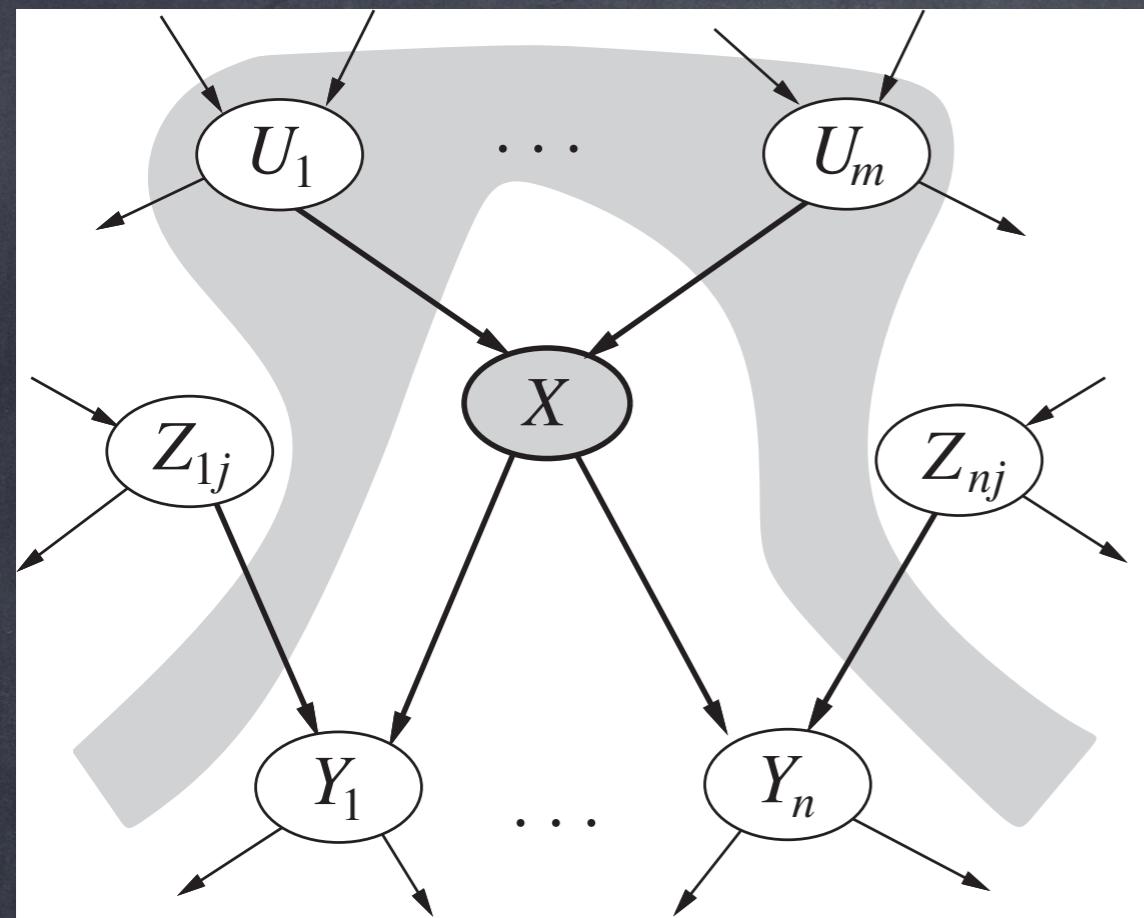
Markov Chain Monte Carlo Simulation

- To approximate: $P(X | e)$
- Generate a sequence of states
 - Values of evidence variables are fixed
 - Values of other variables appear in the right proportion given the distribution encoded by the network

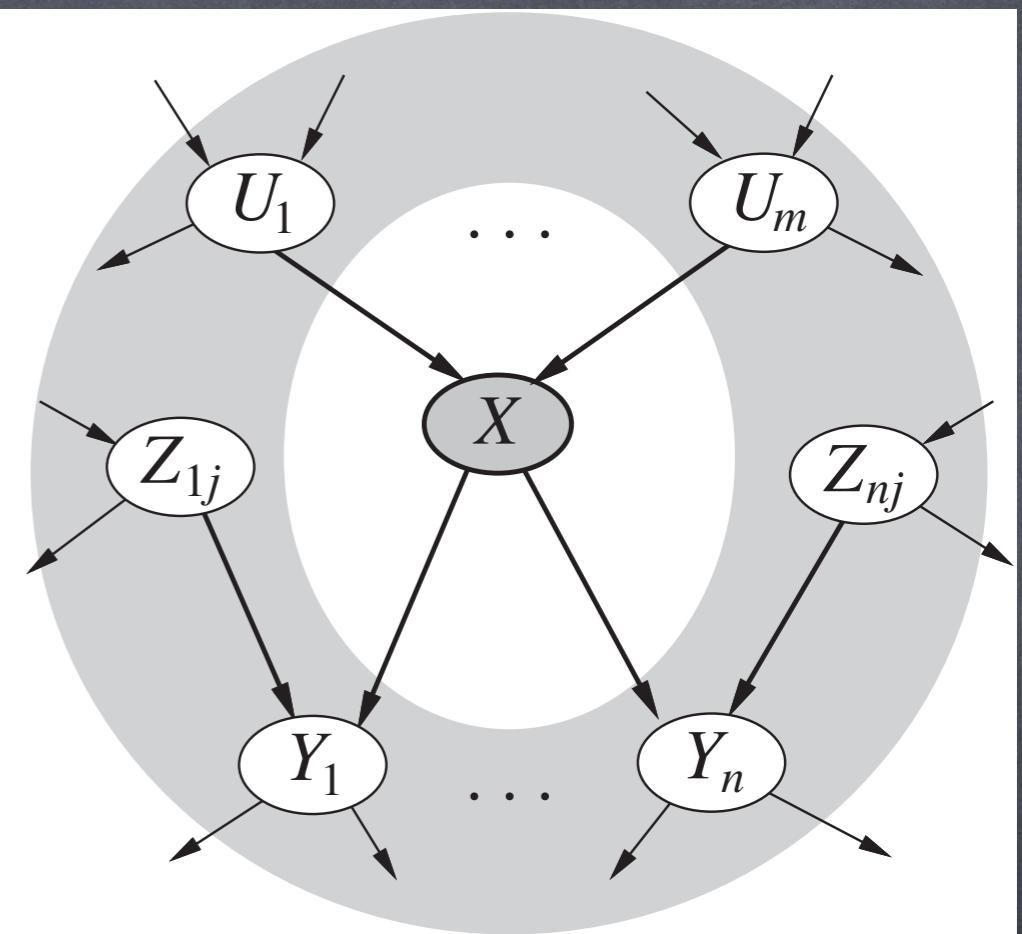


Conditional
Independence

X conditionally independent
of Zs given Us



Conditional
Independence



Markov
Blanket

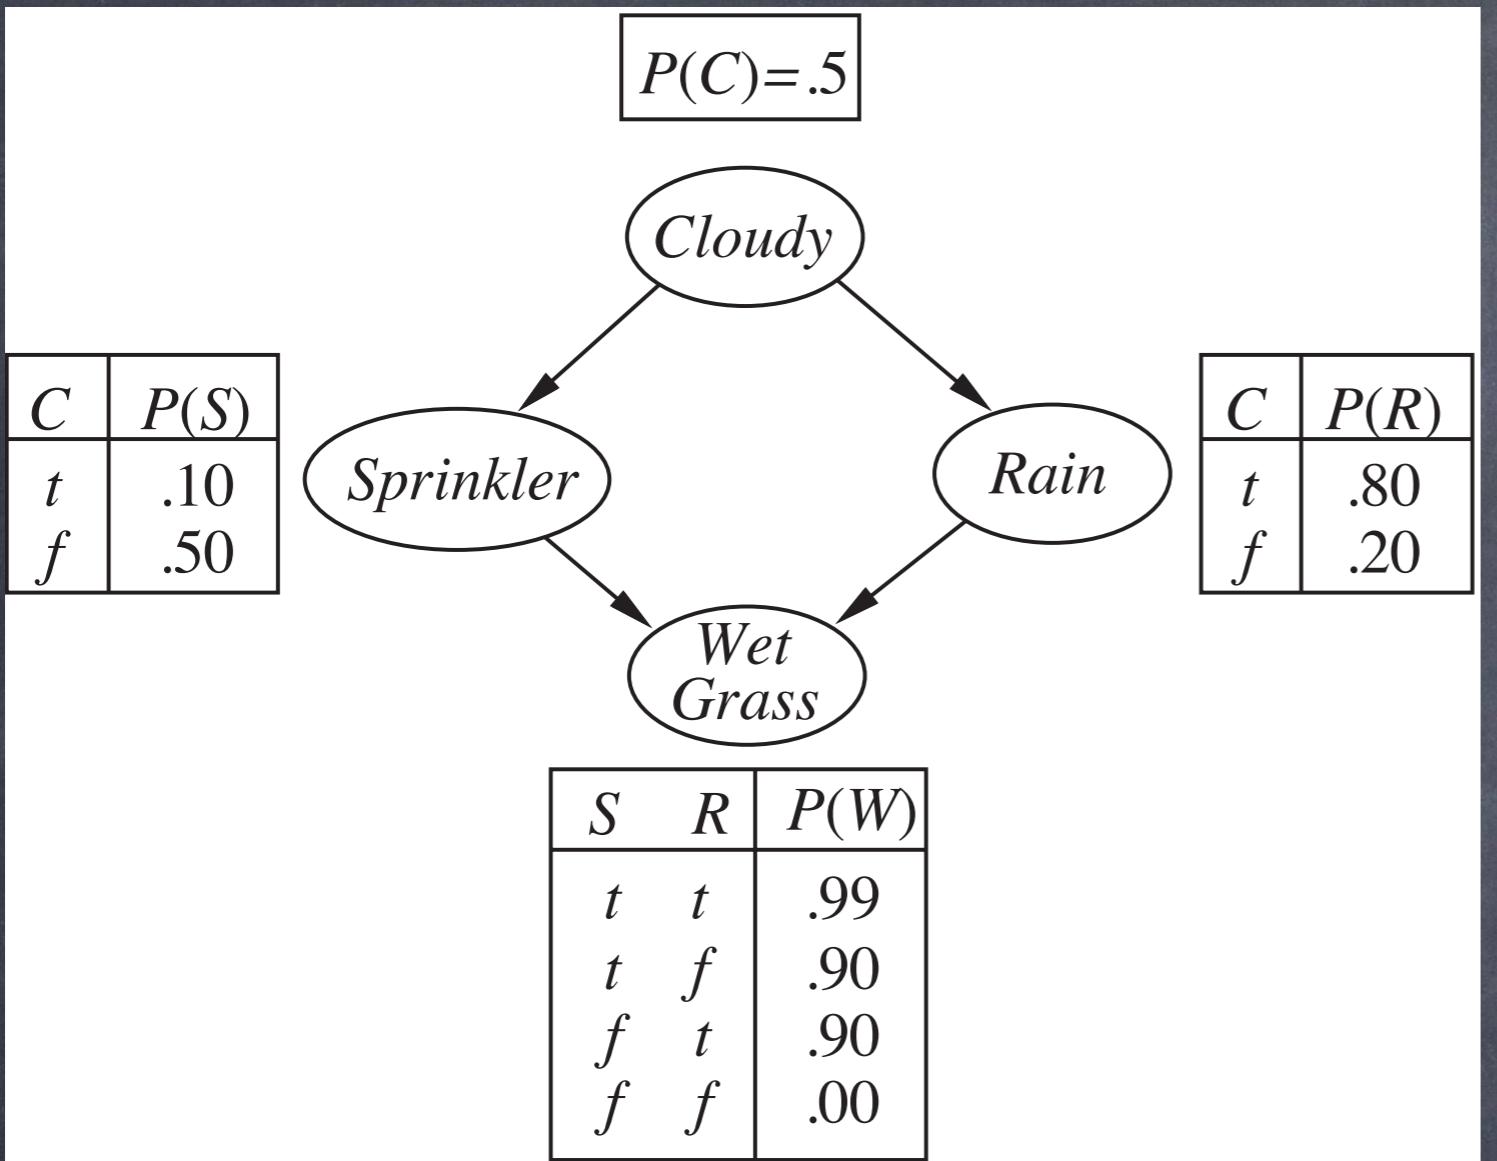
X conditionally independent
of all Vs given Us, Ys, and Zs

Markov Blanket

- The Markov Blanket of a node is its parents, its children, and its children's parents.
- A node is conditionally independent of all other nodes in the network given its Markov Blanket

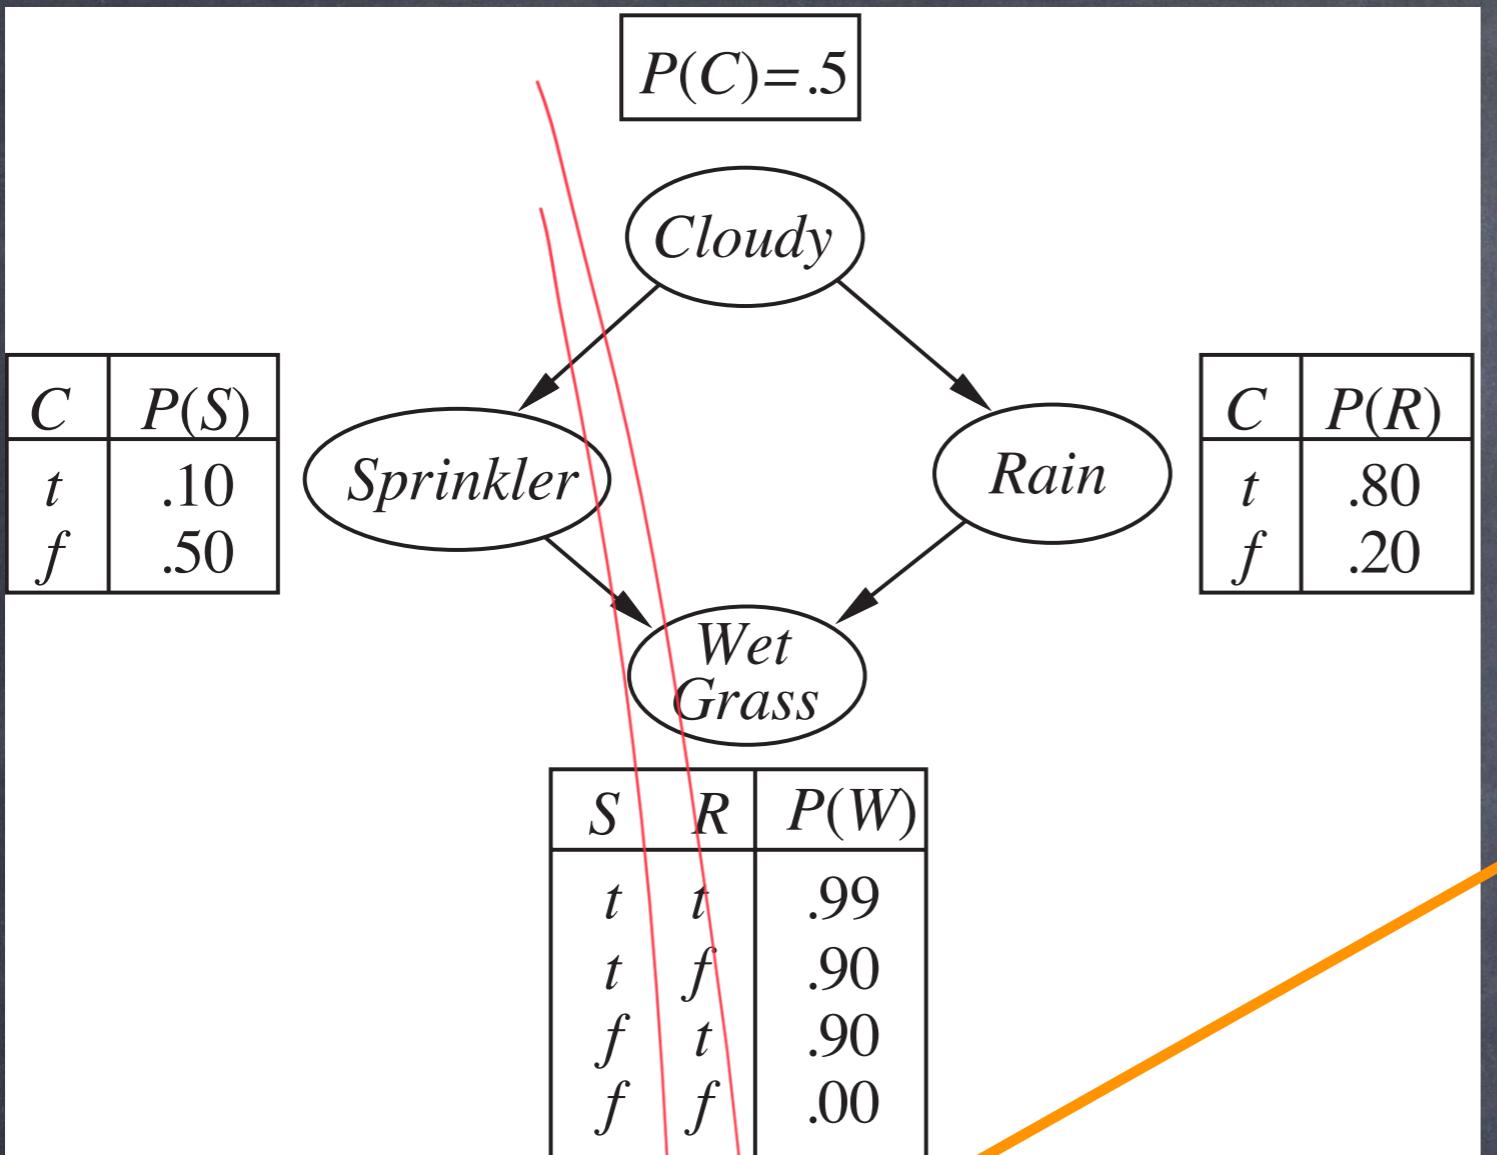
Monte Carlo Simulation

- To approximate: $P(X | e)$
- Start in a state with evidence variables set to evidence values (others arbitrary)
- On each step, sample the non-evidence variables conditioned on the values of the variables in their Markov Blankets
 - Order irrelevant



Cloudy	true
Sprinkler	true
Rain	false
WetGrass	true

$$P(Rain \mid Sprinkler = \text{true}, WetGrass = \text{true})$$

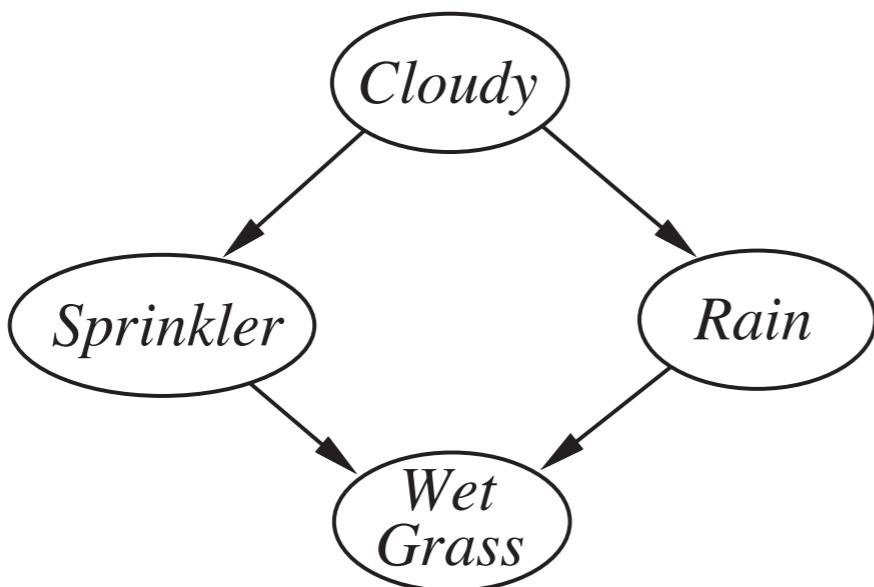


$P(Rain \mid Sprinkler = true, WetGrass = true)$

$N_{Rain=false} += 1$

C	$P(S)$
t	.10
f	.50

$$P(C) = .5$$



C	$P(R)$
t	.80
f	.20

S	R	$P(W)$
t	t	.99
t	f	.90
f	t	.90
f	f	.00

Cloudy

Sprinkler

Rain

WetGrass

true

true

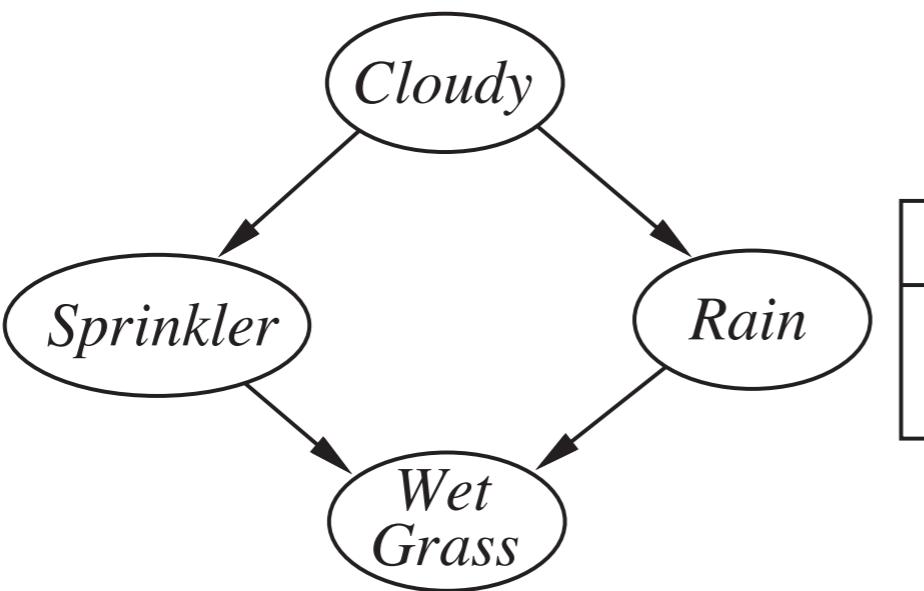
false

true

$$\text{P}(\textit{Rain} \mid \textit{Sprinkler} = \text{true}, \textit{WetGrass} = \text{true})$$

C	$P(S)$
t	.10
f	.50

$$P(C) = .5$$



C	$P(R)$
t	.80
f	.20

S	R	$P(W)$
t	t	.99
t	f	.90
f	t	.90
f	f	.00

Cloudy

Sprinkler

Rain

WetGrass

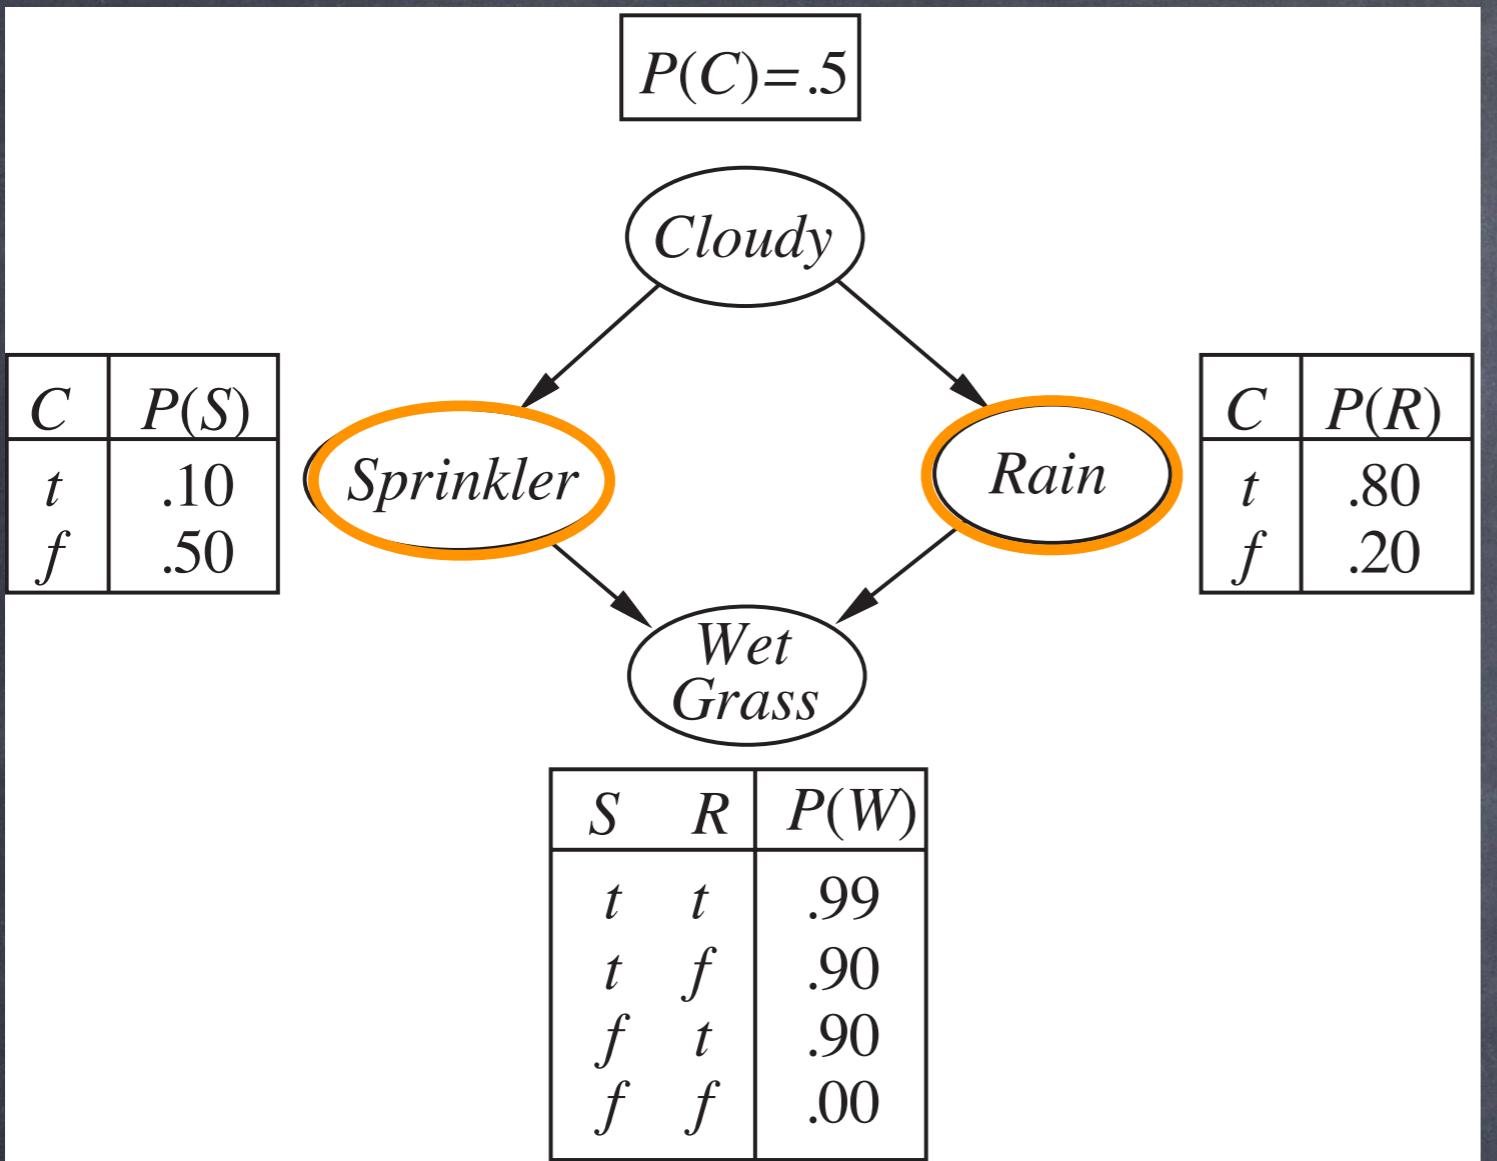
true

true

false

true

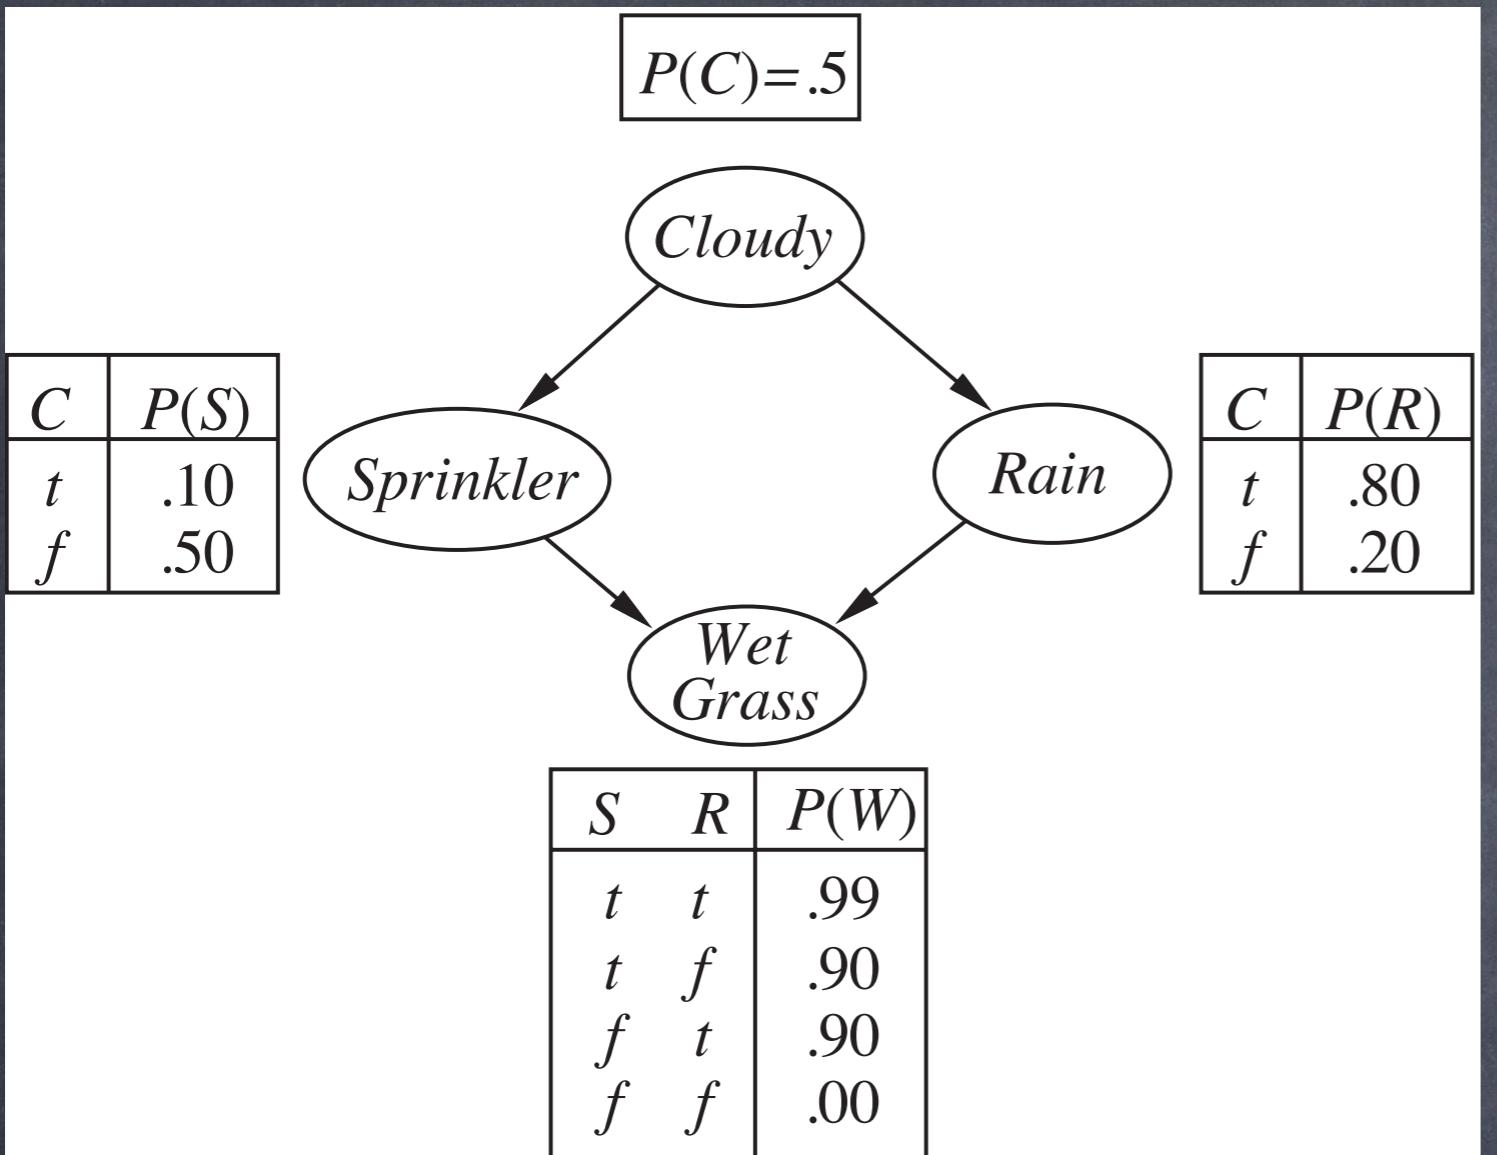
$$\mathbf{P}(Rain \mid Sprinkler = true, WetGrass = true)$$



Cloudy	true
Sprinkler	true
Rain	false
WetGrass	true

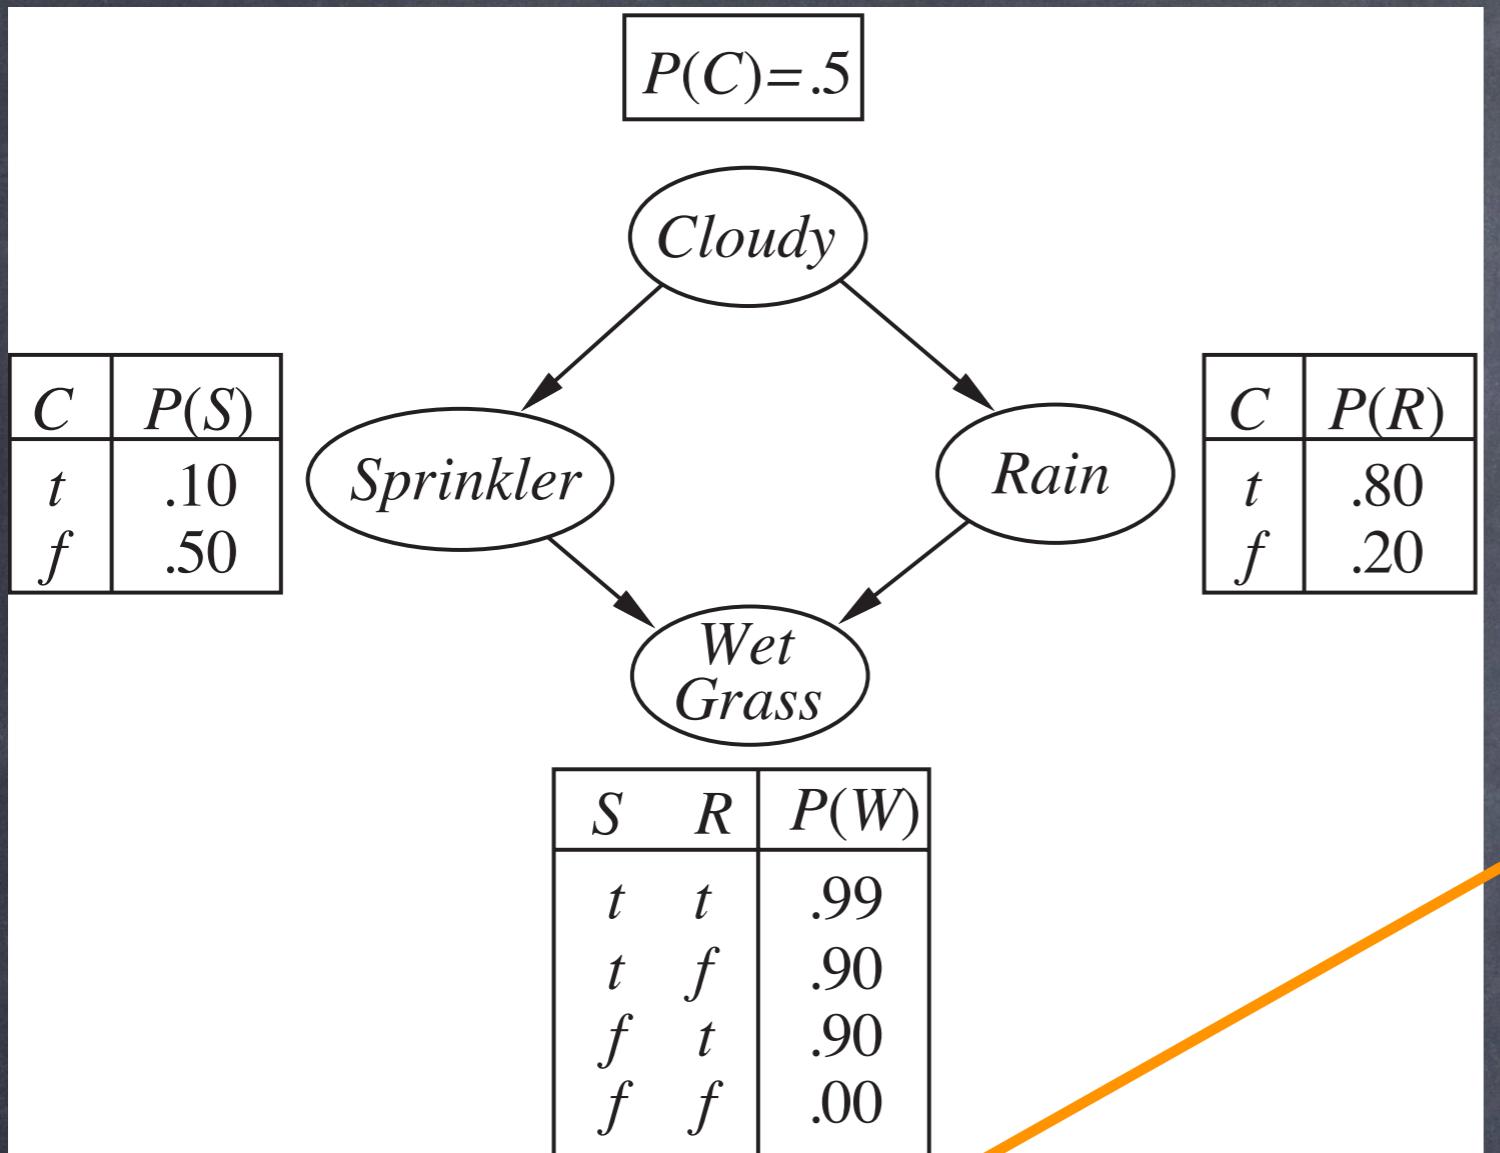
$$P(Rain \mid Sprinkler = \text{true}, WetGrass = \text{true})$$

$$P(\text{Cloudy} \mid \text{Sprinkler}=\text{true}, \text{Rain}=\text{false})$$



Cloudy	false
Sprinkler	true
Rain	false
WetGrass	true

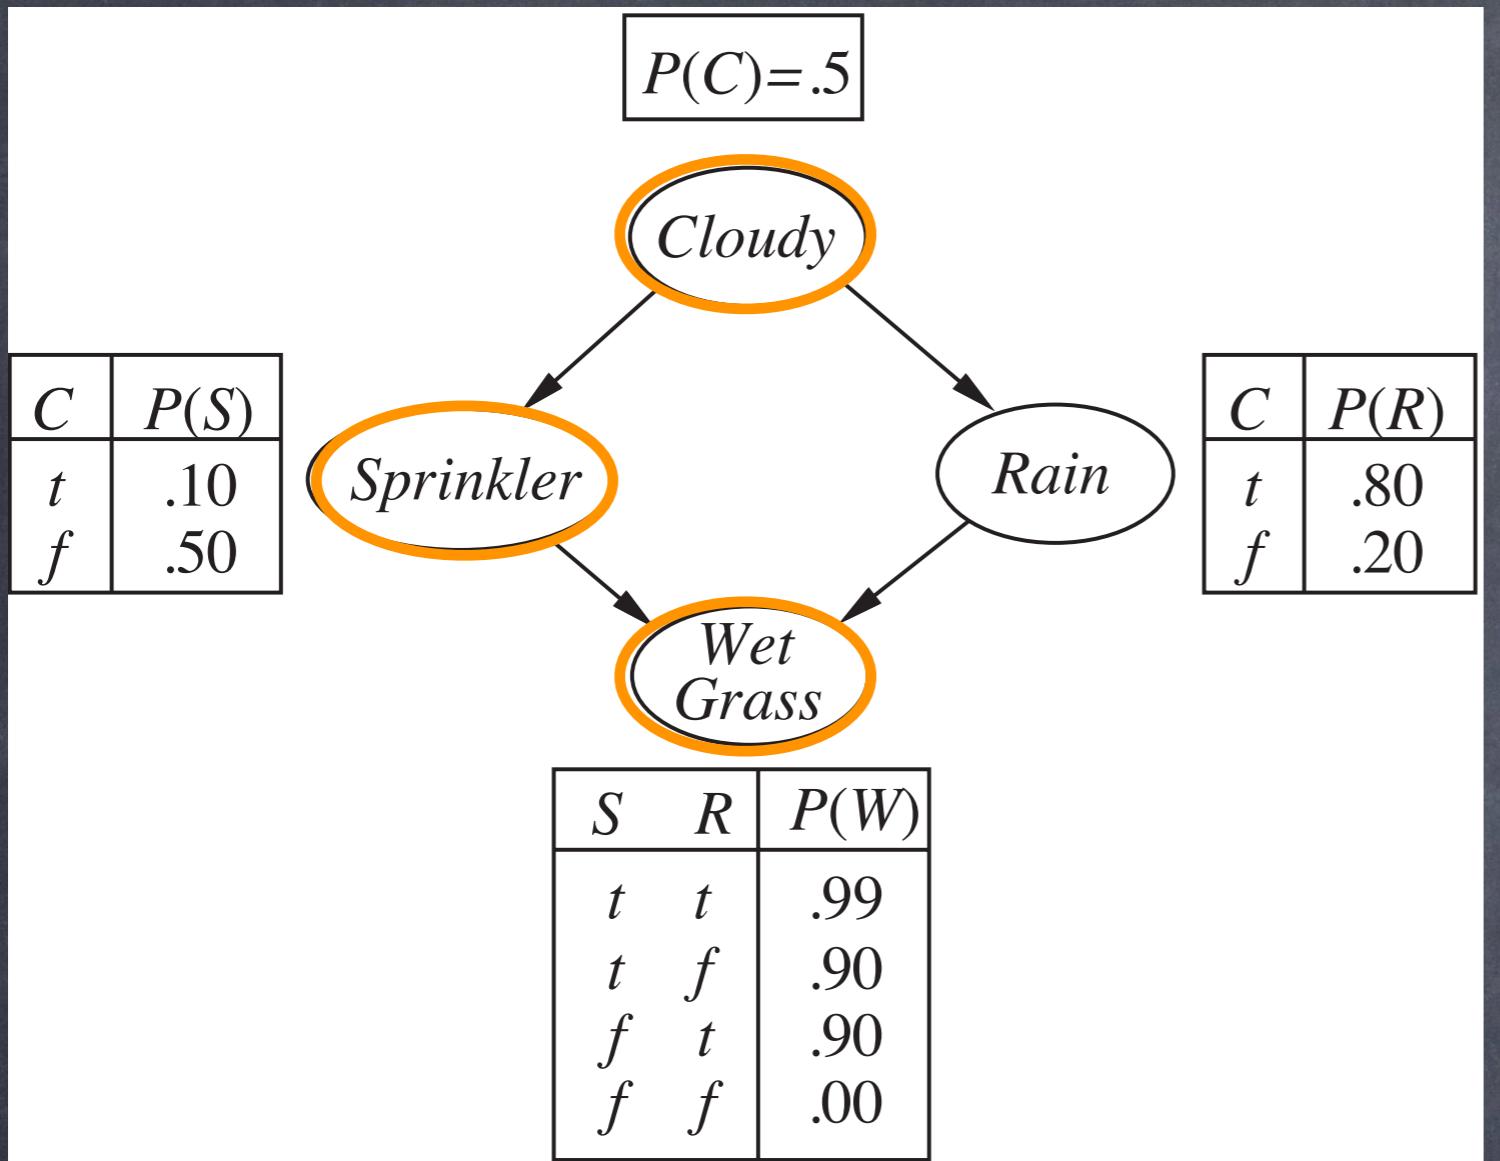
$$P(Rain \mid Sprinkler = \text{true}, WetGrass = \text{true})$$



Cloudy	false
Sprinkler	true
Rain	false
WetGrass	true

$$P(Rain \mid Sprinkler = \text{true}, WetGrass = \text{true})$$

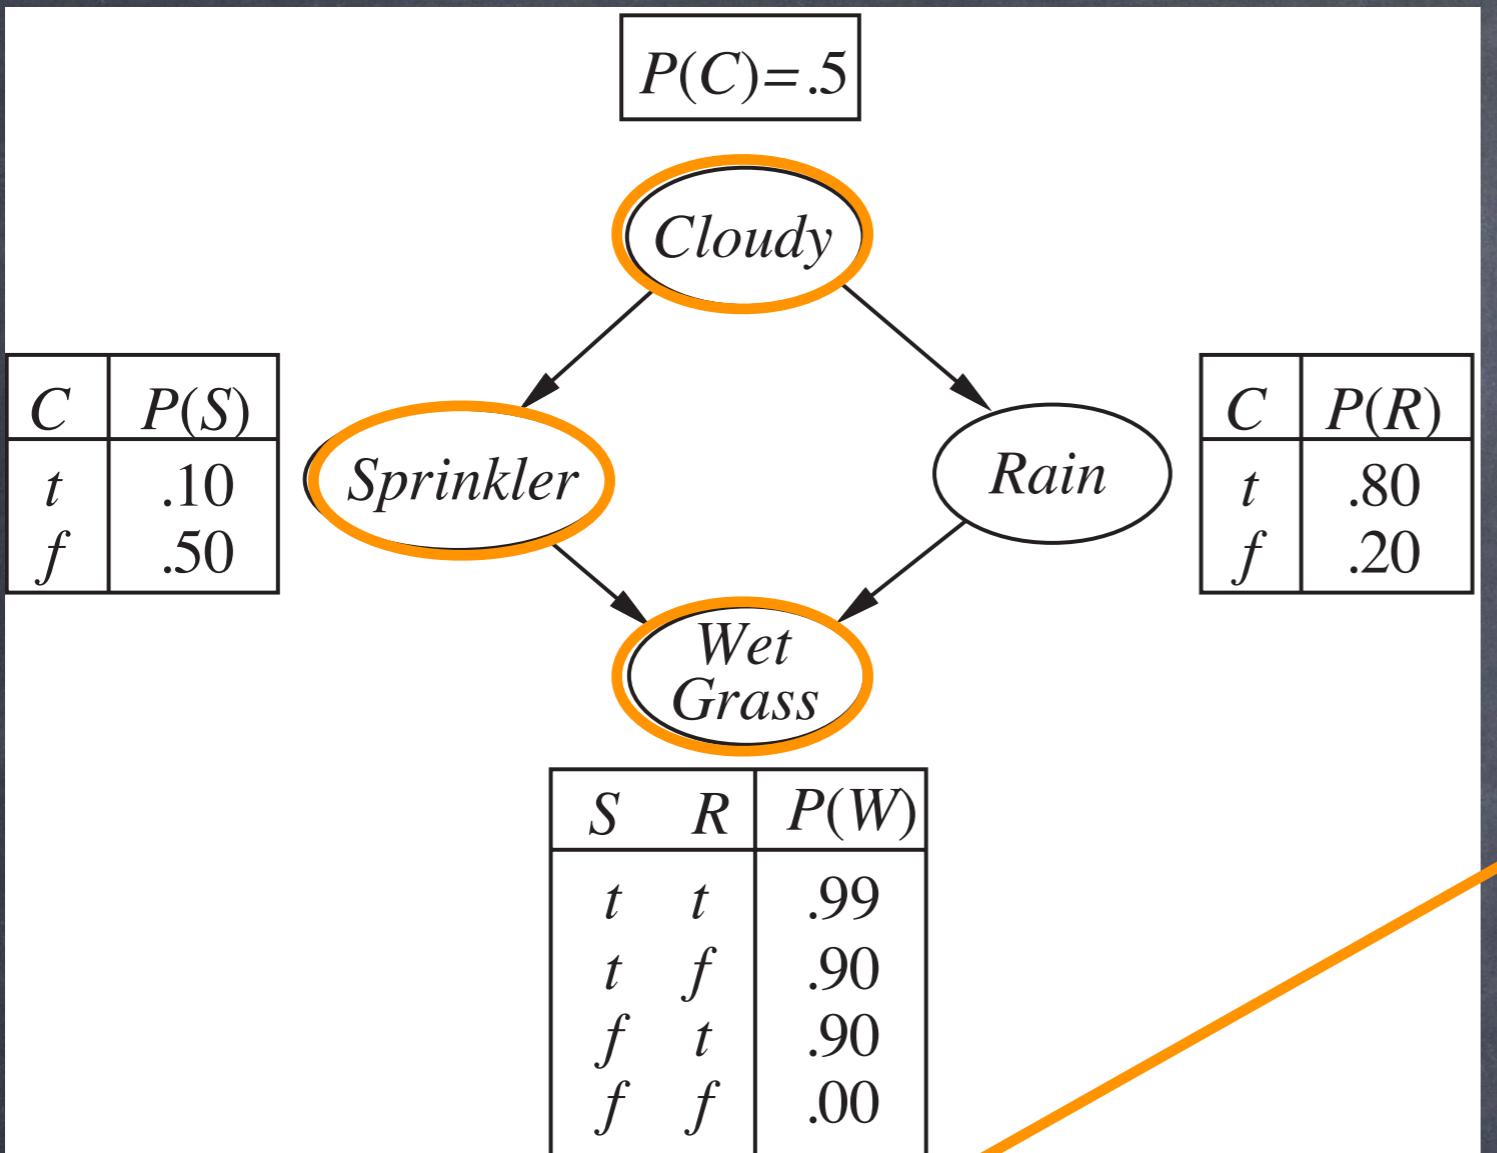
$N_{Rain=false} += 1$



Cloudy	false
Sprinkler	true
Rain	false
WetGrass	true

$$P(Rain \mid Sprinkler = true, WetGrass = true)$$

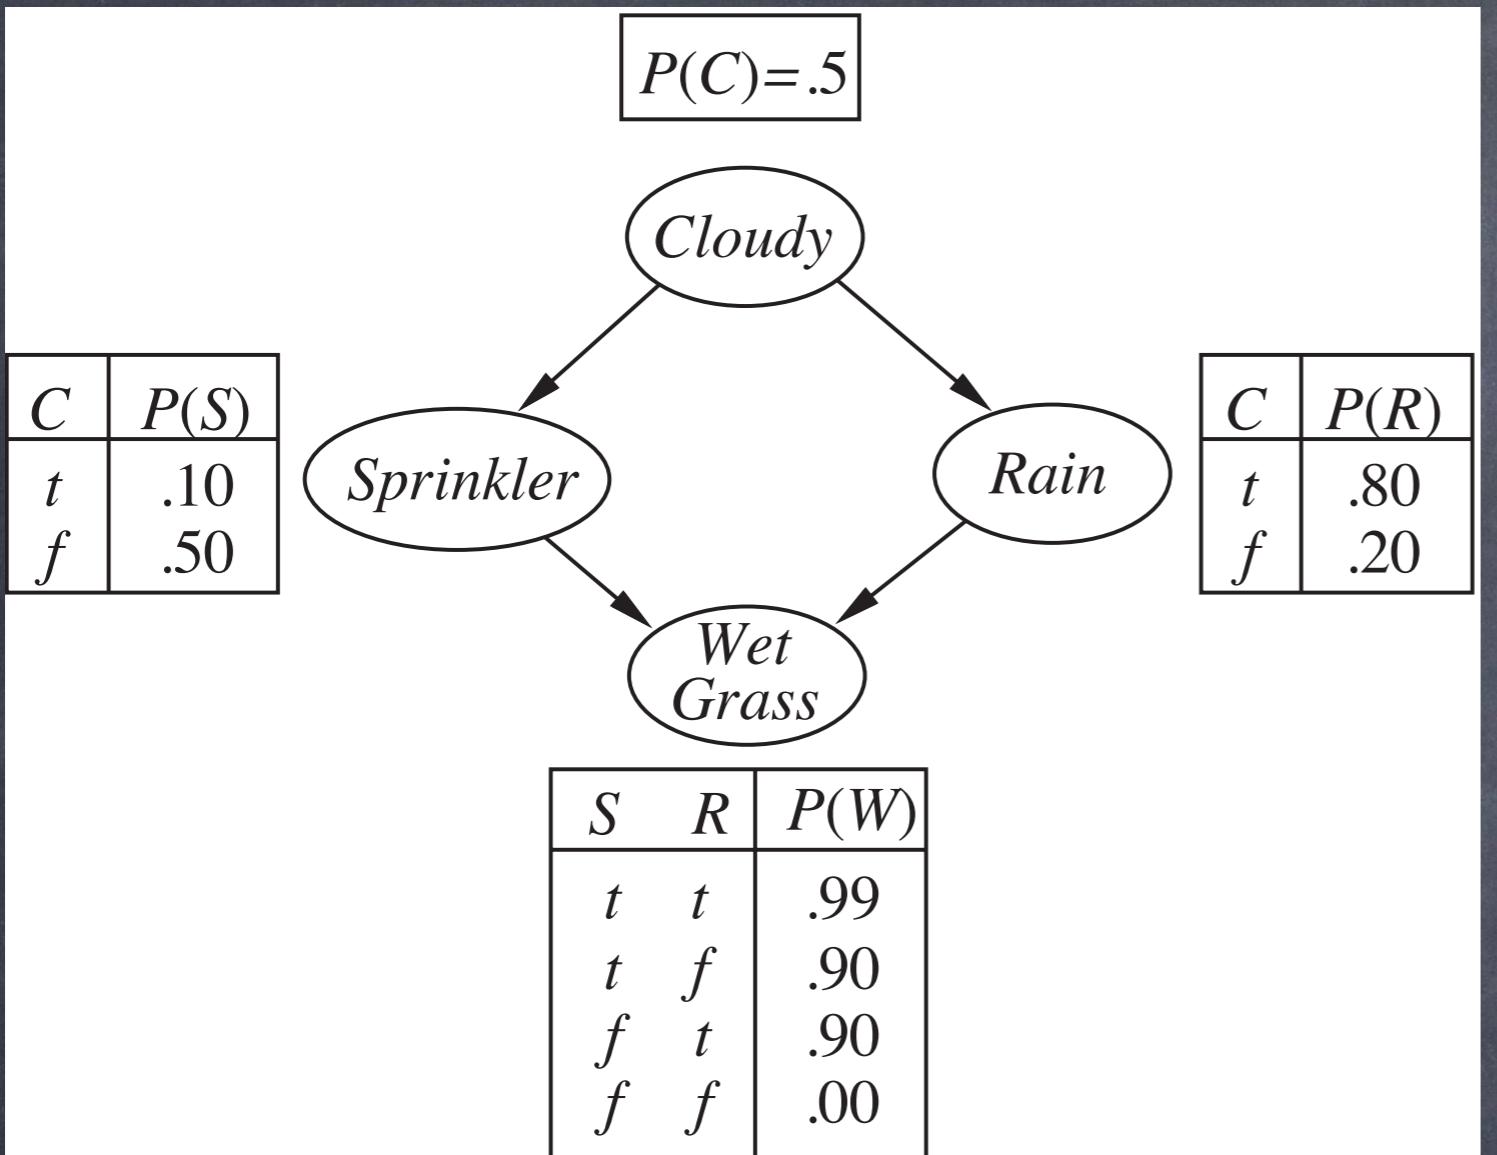
P(Rain | Sprinkler=true, Rain=false, Cloudy=false)



Cloudy	false
Sprinkler	true
Rain	true
WetGrass	true

$$\mathbf{P}(Rain \mid Sprinkler = true, WetGrass = true)$$

`NRain=true += 1`



Cloudy	false
Sprinkler	true
Rain	true
WetGrass	true

$$P(Rain \mid Sprinkler = \text{true}, WetGrass = \text{true})$$

Gibbs Sampling

- To approximate: $P(X | e)$
- Start in a state with evidence variables set to evidence values (others arbitrary)
- On each step, sample non-evidence variables conditioned on the values of the variables in their Markov Blanket
 - Order irrelevant
 - A form of local search!

Exact Inference in Bayesian Networks

$$\begin{aligned} \mathbf{P}(X \mid \mathbf{e}) &= \alpha \mathbf{P}(X, \mathbf{e}) = \alpha \sum_{\mathbf{y}} \mathbf{P}(X, \mathbf{e}, \mathbf{y}) \\ &= \alpha \sum_{\mathbf{y}} \prod_{i=1}^n P(X_i \mid \text{parents}(X_i)) \end{aligned}$$

- Intractable (#P-Hard)

Approximate Inference in Bayesian Networks

- Sampling consistent with a distribution
- Rejection Sampling: rejects too much
- Likelihood Weighting: weights get too small
- Gibbs Sampling: MCMC algorithm
 - Similar to local search
- All generate consistent estimates (equal to exact probability in the large-sample limit)